

# Assignment - 2

## Q1 N-Cuts

- **N-Cut for segmenting brightness Images**

- Feature vector of a pixel used for calculating weight

$$F(i) = (I_{\text{red}}(i) + I_{\text{green}}(i) + I_{\text{blue}}(i)) / 3$$

$I_{\text{red}}(i)$  -> Intensity of pixel in red channel

$I_{\text{green}}(i)$  -> Intensity of pixel in green channel

$I_{\text{blue}}(i)$  -> Intensity of pixel in blue channel

- Each element ( $W[i][j]$ ) in the weight matrix is calculated using feature proximity and spatial proximity terms.

$$\text{Spatial\_proximity}[i][j] = \exp(-(\text{distance}(i, j)^2) / \sigma_x^2)$$

$$\text{Feature\_proximity}[i][j] = \exp(-(|F(i) - F(j)|^2) / \sigma_i^2)$$

```
if distance(i, j) < r
    W[i][j] = Feature_proximity[i][j] * (Spatial_proximity[i][j])
```

else

$$W[i][j] = 0$$

r-> cutoff distance

- Calculated degree matrix (D) as follows

$$D = np.diag(W.sum(axis=1))$$

- Calculated eigen vector corresponding to the second smallest eigenvalue for the following generalised eigen system-

$$(D - W) * y = \lambda * D * y$$

- Used threshold = 0 for eigen vector

- Parameters used-

$$R = 50$$

$$\Sigma_i = 10$$

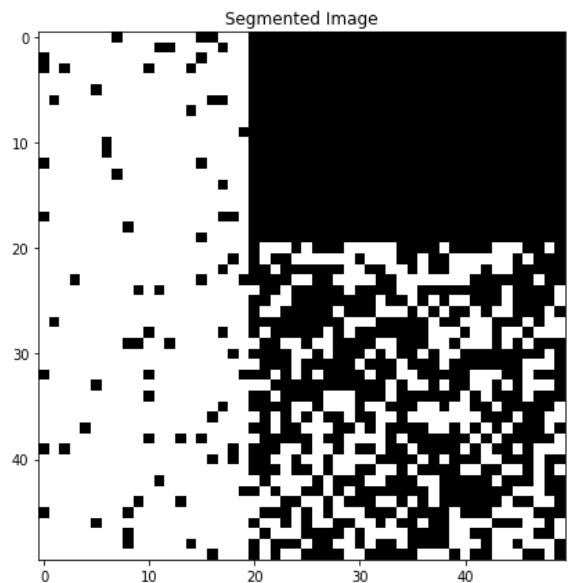
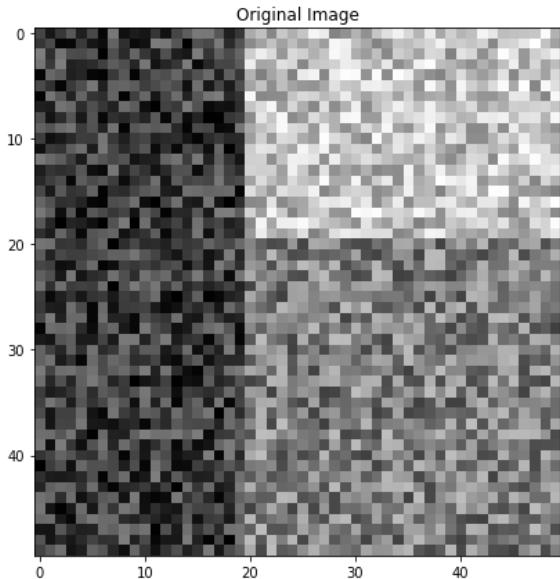
$$\Sigma_x = 1e7$$

- Downscaled the images to 50\*50 for faster computation

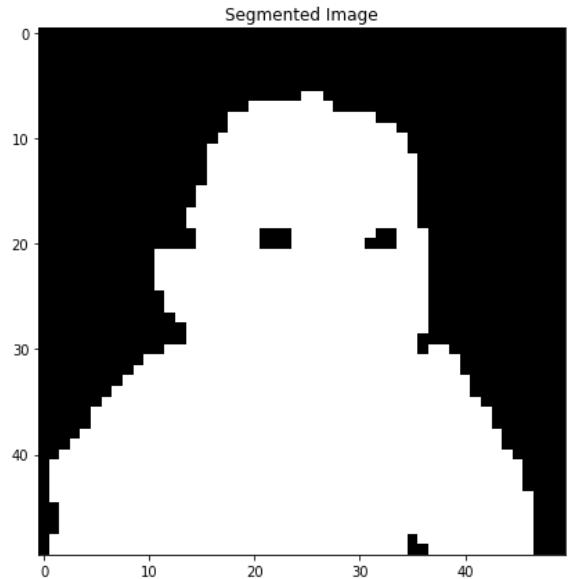
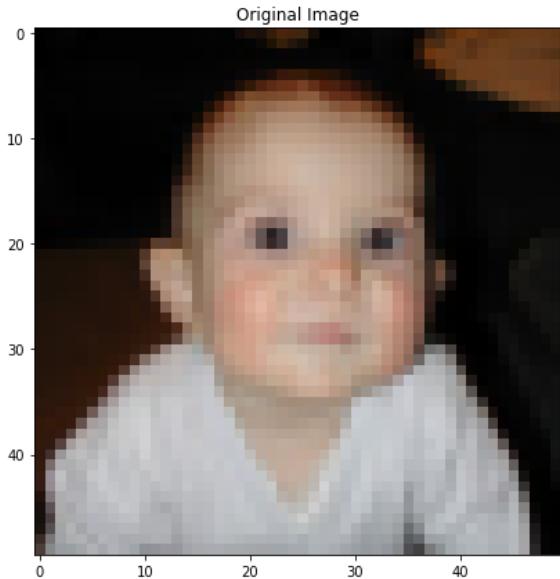
## **Results:**

### i) N-Cut on Original Image (Downscaled)

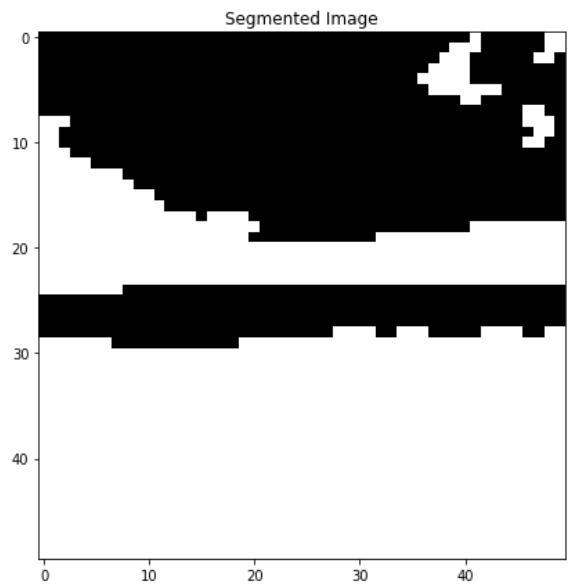
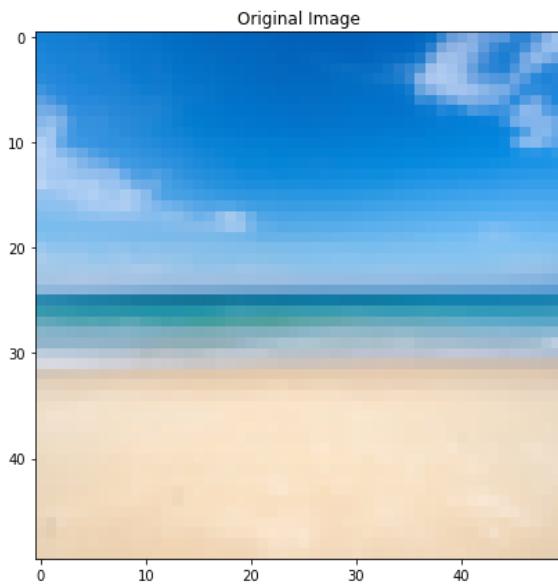
test1.png (downscaled - (50\*50) )



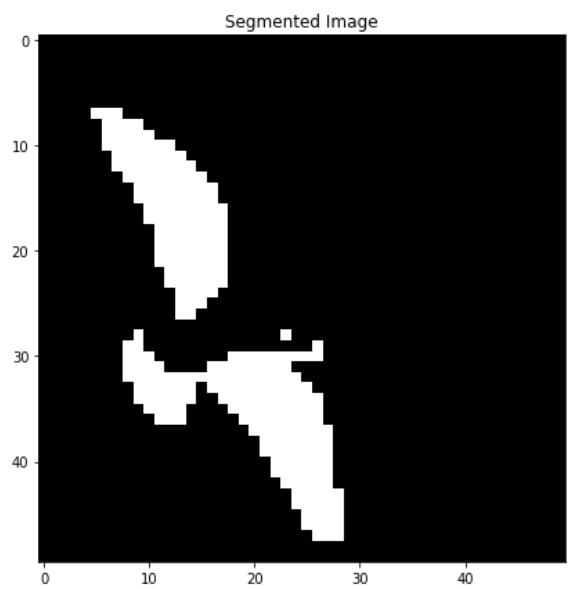
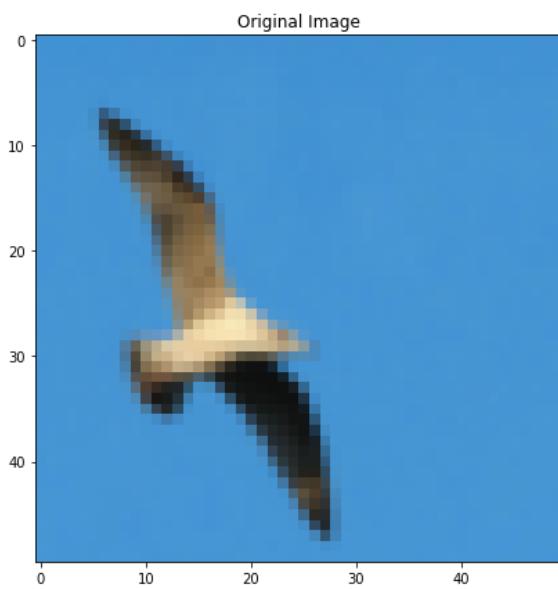
test2.jpg (downscaled - (50\*50) )



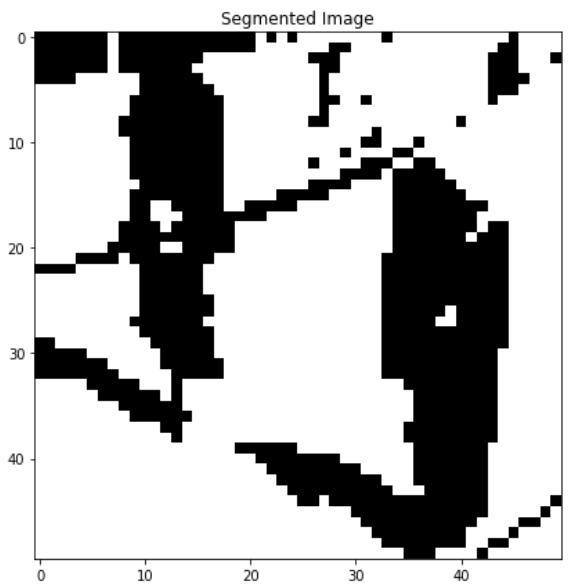
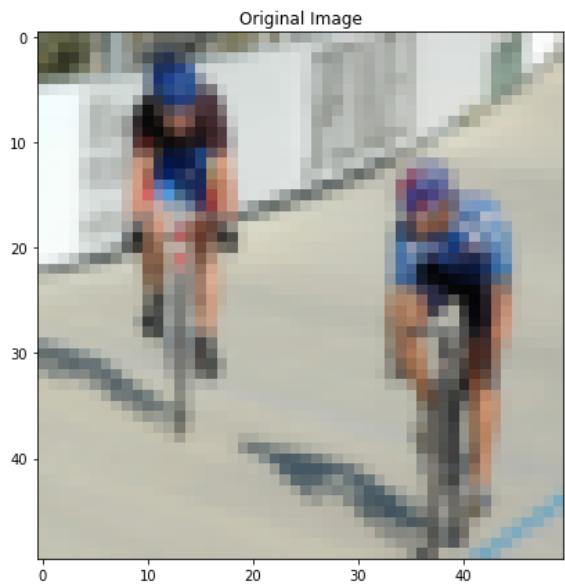
test3.jpg (downscaled - (50\*50)



test4.jpg (downscaled - (50\*50)

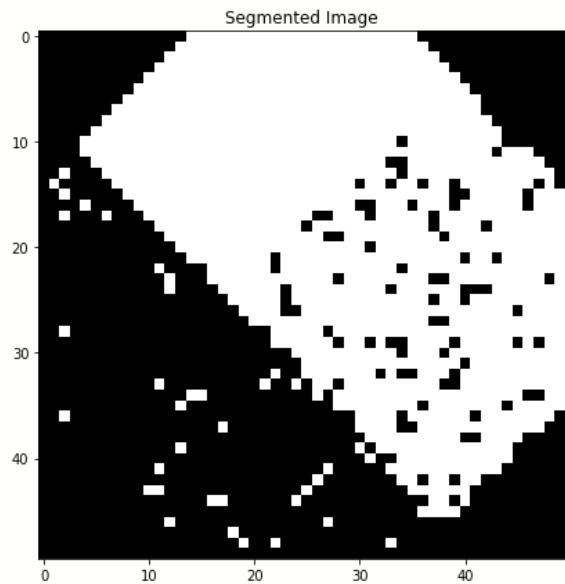
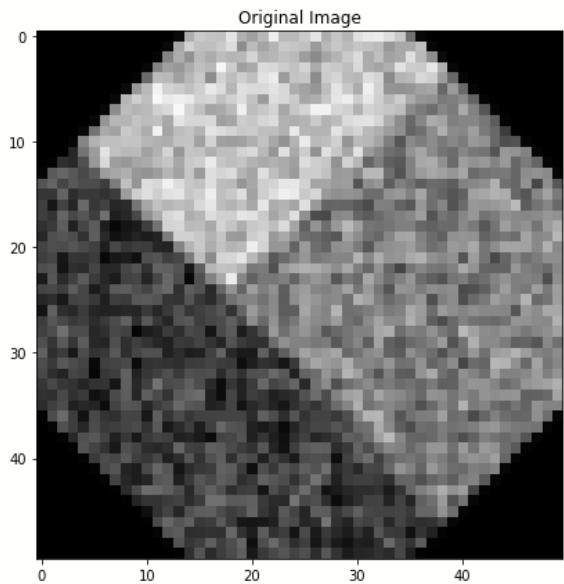


test5.jpg (downscaled - (50\*50) )

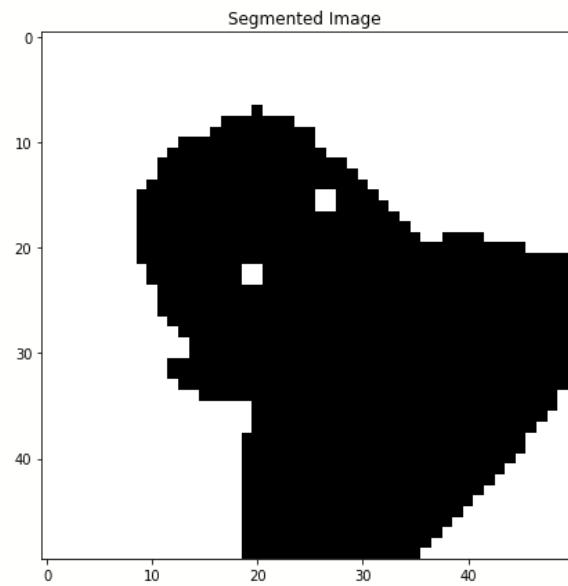
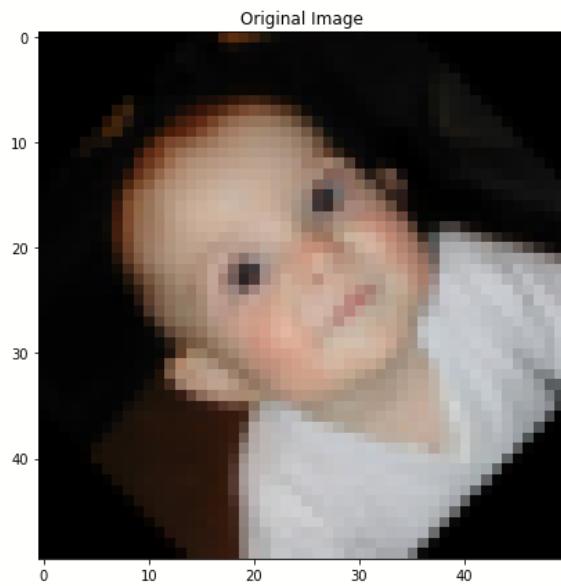


ii) N-Cut on Rotated Image (Downscaled)

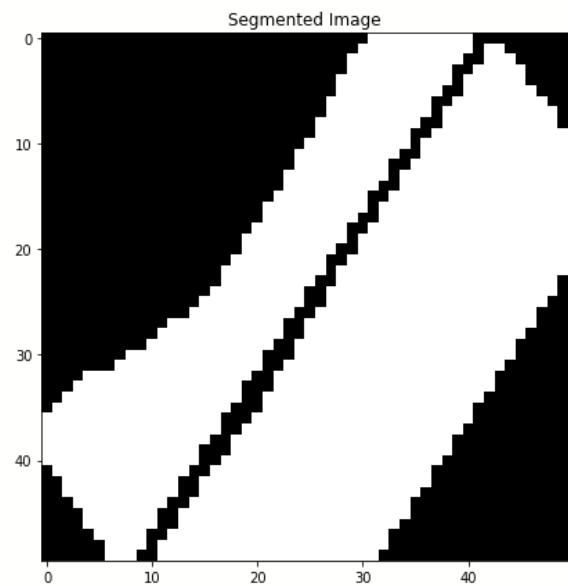
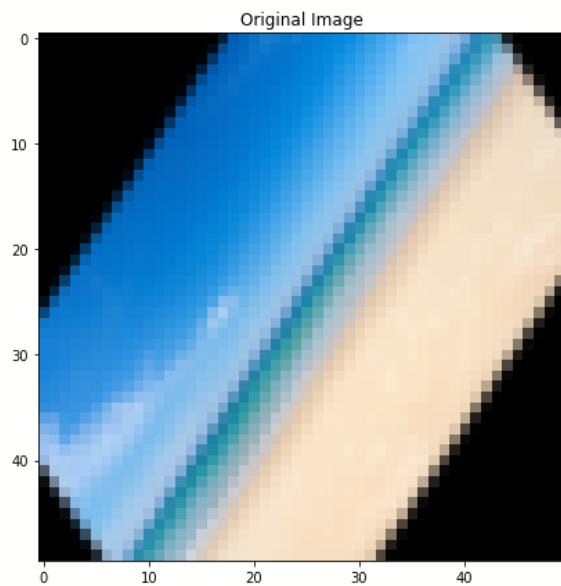
test1.png (downscaled - (50\*50) )



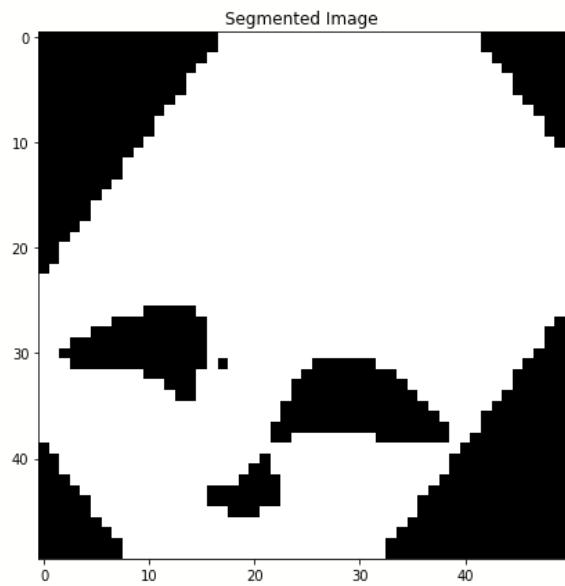
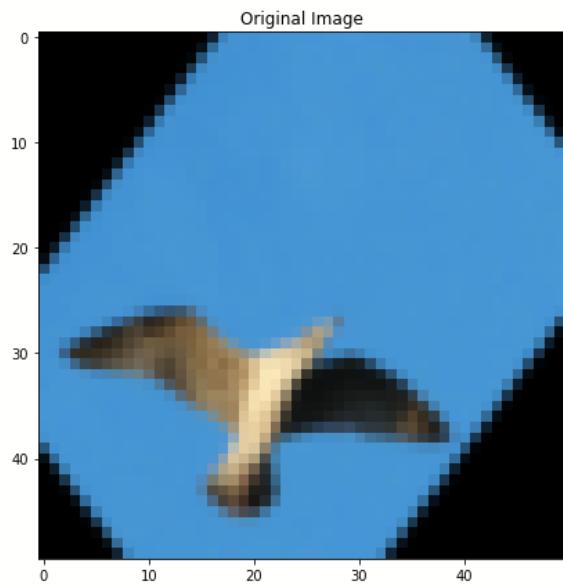
test2.jpg (downscaled - (50\*50))



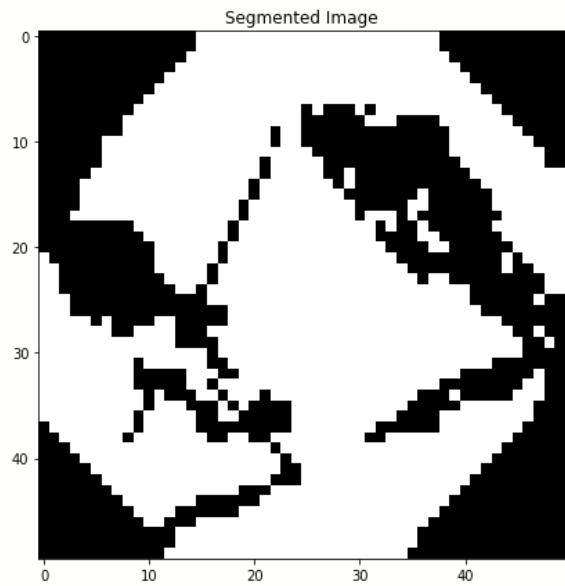
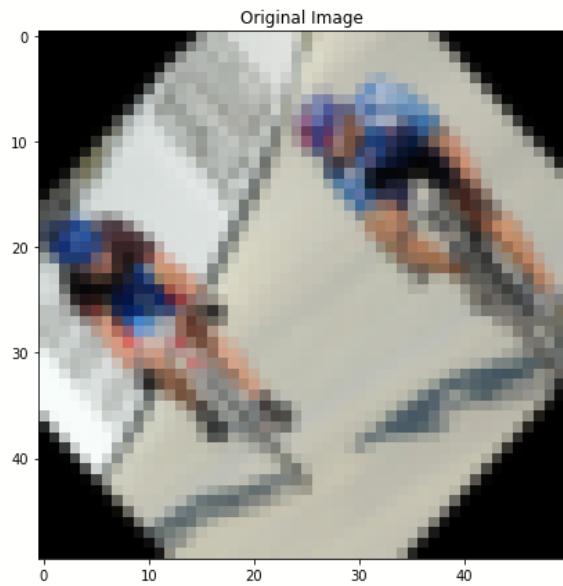
test3.jpg (downscaled - (50\*50) )



test4.jpg (downscaled - (50\*50) )

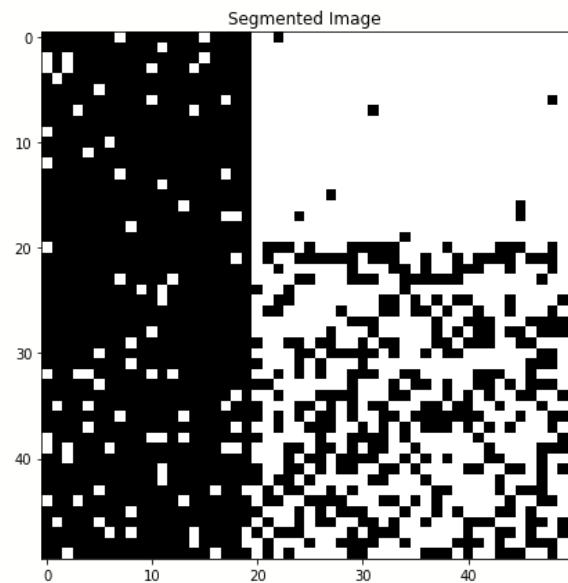
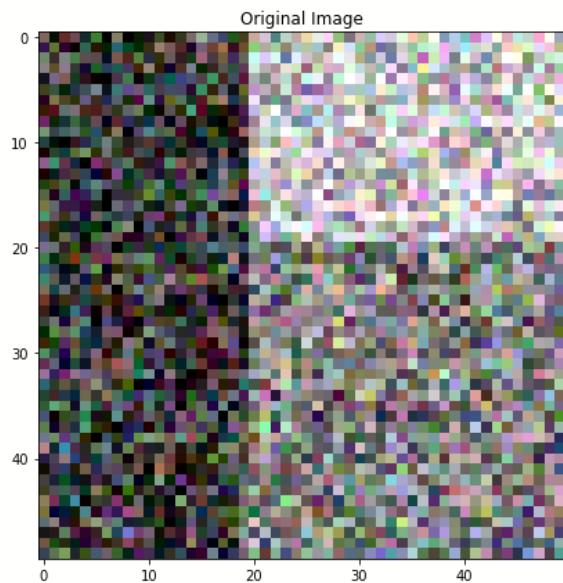


test5.jpg (downscaled - (50\*50) )

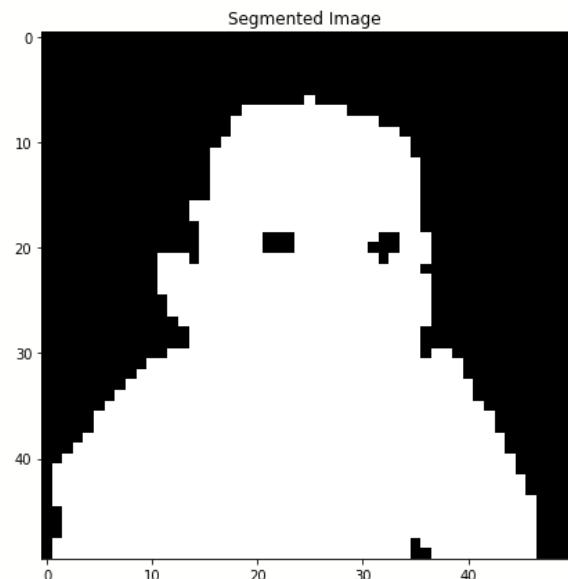
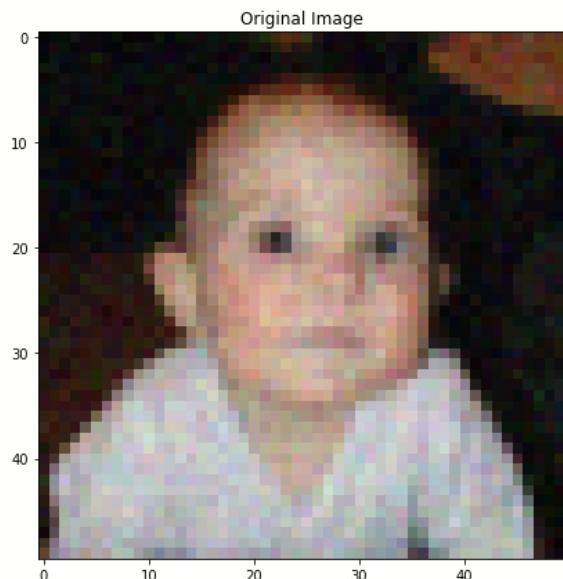


**iii) N-Cut on Gaussian Blurred Image (Downscaled)**

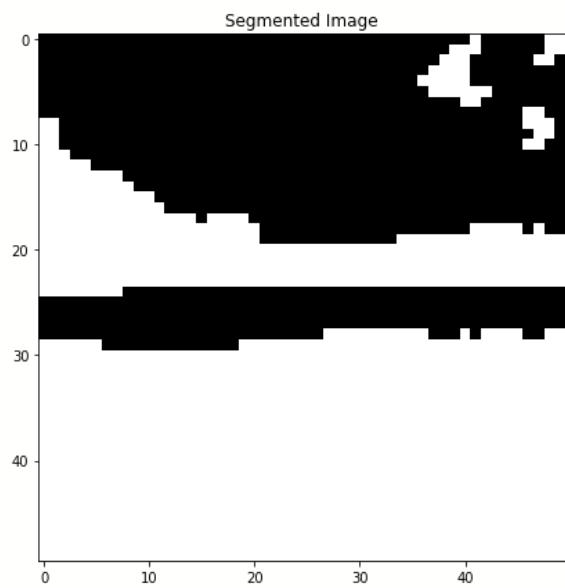
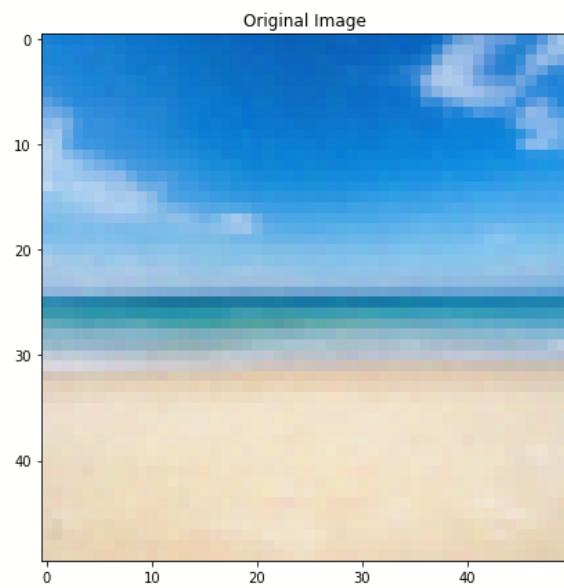
test1.png (downscaled - (50\*50) )



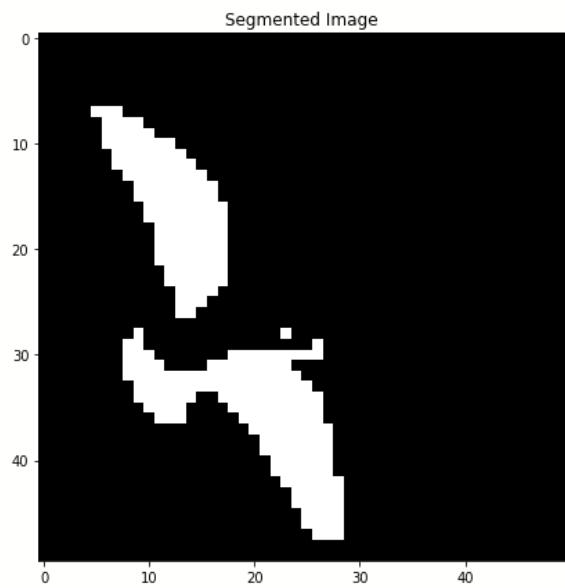
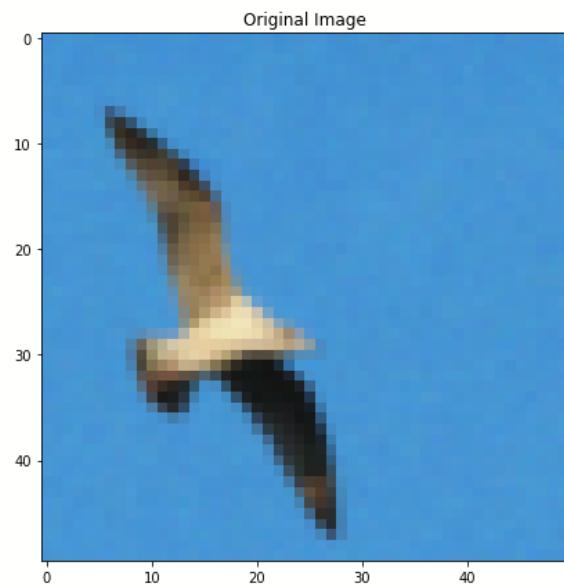
test2.jpg (downscaled - (50\*50) )



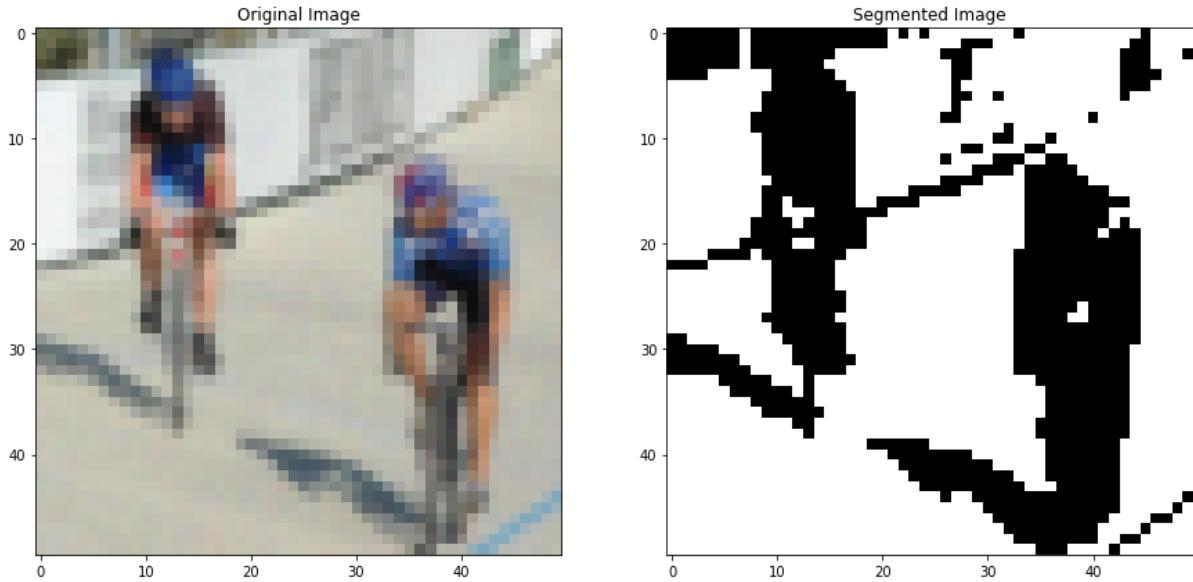
test3.jpg (downscaled - (50\*50) )



test4.jpg (downscaled - (50\*50) )



test5.jpg (downscaled - (50\*50) )



- **N-Cut for segmenting color Images**

- Feature vector of a pixel used for calculating weight

$$F(i) = [v, v*s*\sin(h), v*s*\cos(h)]$$

h,s,v are hsv values of the pixel

- Each element ( $W[i][j]$ ) in the weight matrix is calculated using feature proximity and spatial proximity terms.

$$\text{Spatial\_proximity } [i][j] = \exp(-(\text{distance}(i, j)^2) / \sigma_x^2)$$

$$\text{Feature\_proximity } [i][j] = \exp(-(\text{norm}(F(i) - F(j))^2) / \sigma_i^2)$$

```
if distance(i, j) < r
    W[i][j] = Feature_proximity [i][j] * (Spatial_proximity [i][j])
else
    W[i][j] = 0
```

r-> cutoff distance

- Calculated degree matrix (D) as follows

$$D = \text{np.diag}(W.\text{sum}(\text{axis} = 1))$$

- Calculated eigen vector corresponding to the second smallest eigenvalue for the following generalised eigen system-

$$(D-W) * y = \lambda * D * y$$

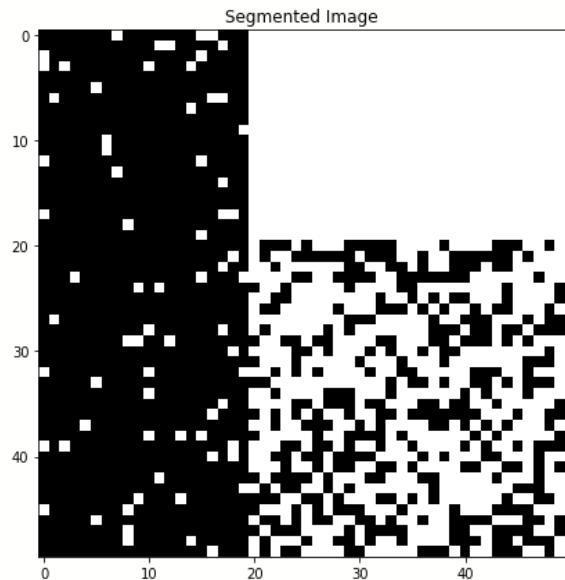
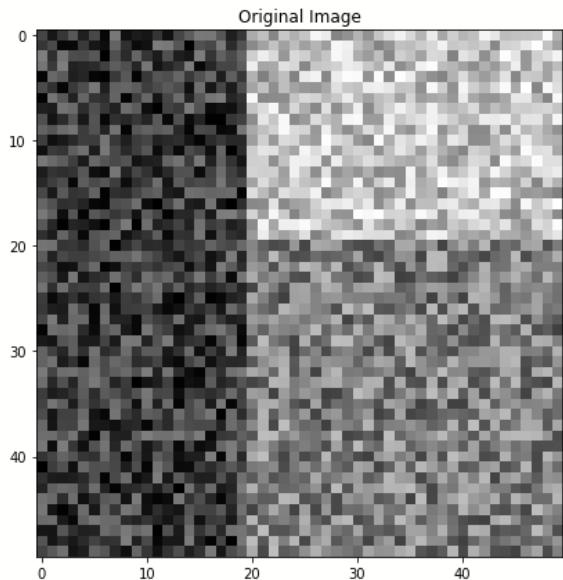
- Used threshold = 0 for eigen vector

- Parameters used-
   
R = 50
   
Sigma\_i = 10
   
Sigma\_x = 1e7
- Downscaled the images to 50\*50 for faster computation

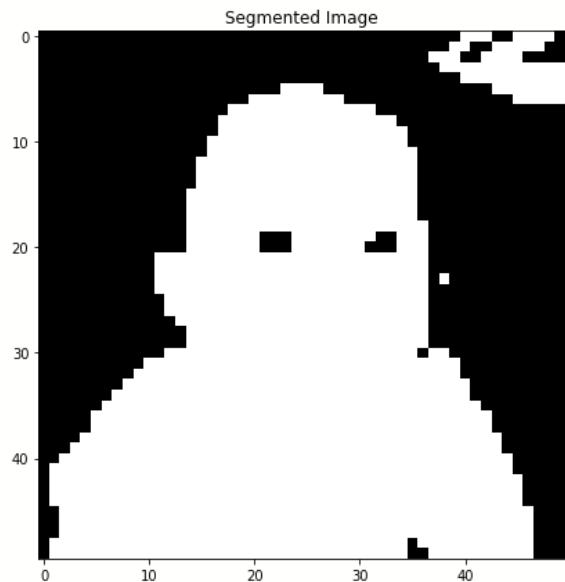
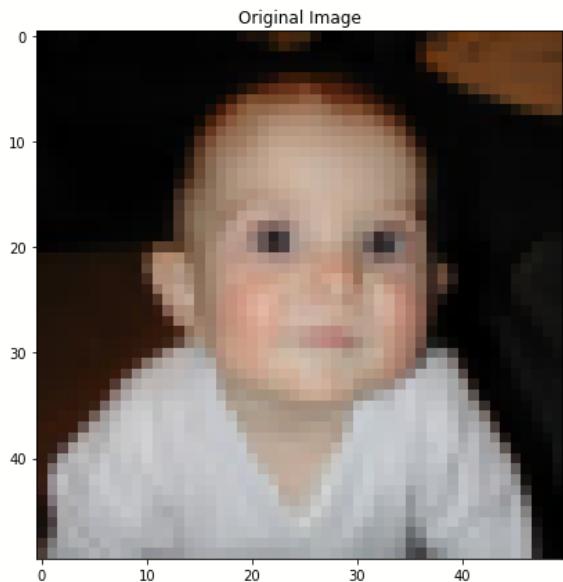
## Results:

### i) N-Cut on Original Image (Downscaled)

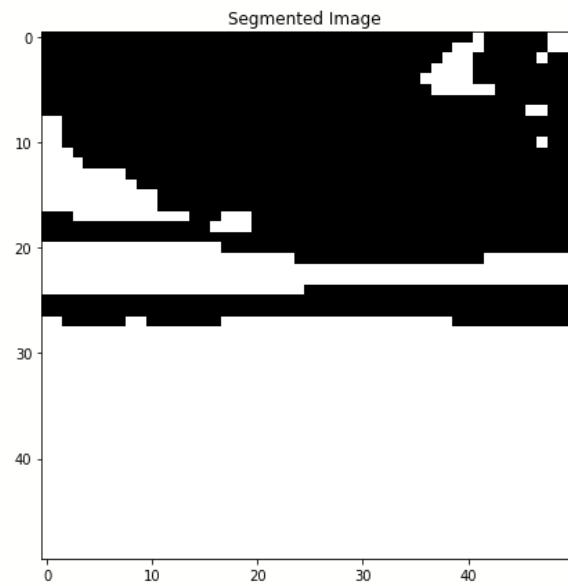
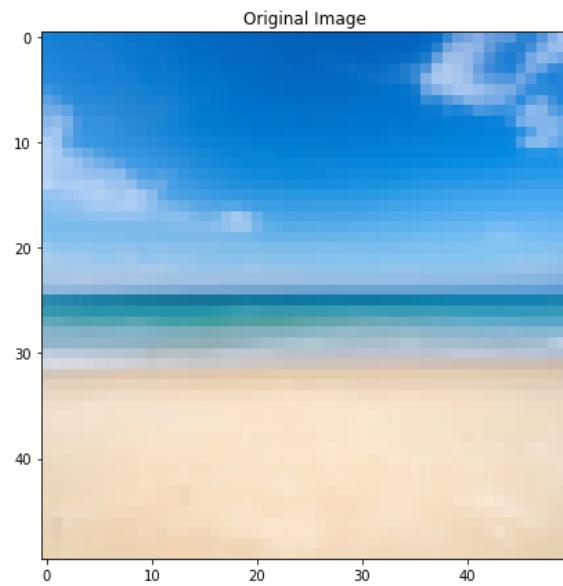
test1.png (downscaled - (50\*50) )



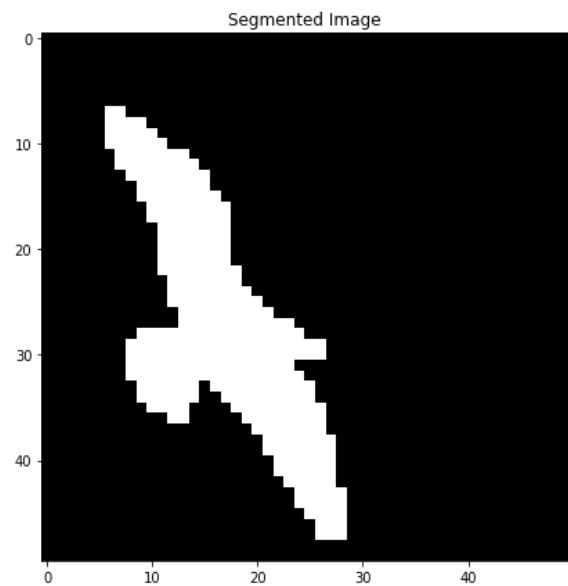
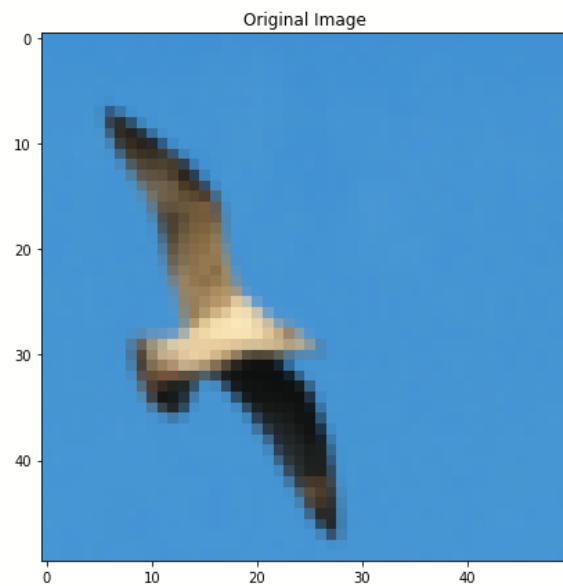
test2.jpg (downscaled - (50\*50) )



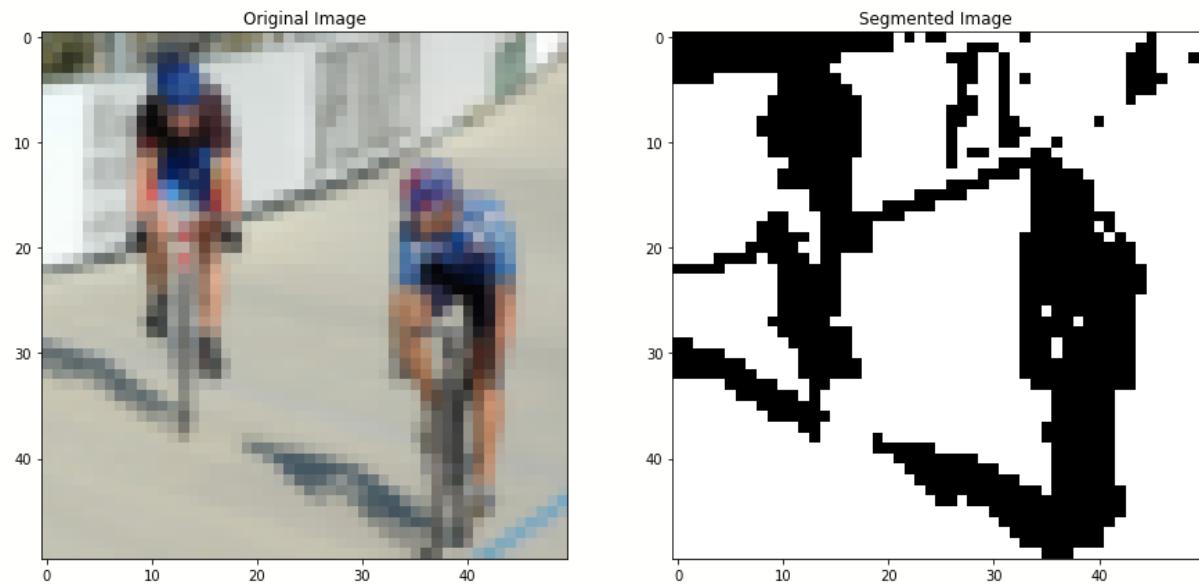
test3.jpg (downscaled - (50\*50) )



test4.jpg (downscaled - (50\*50) )

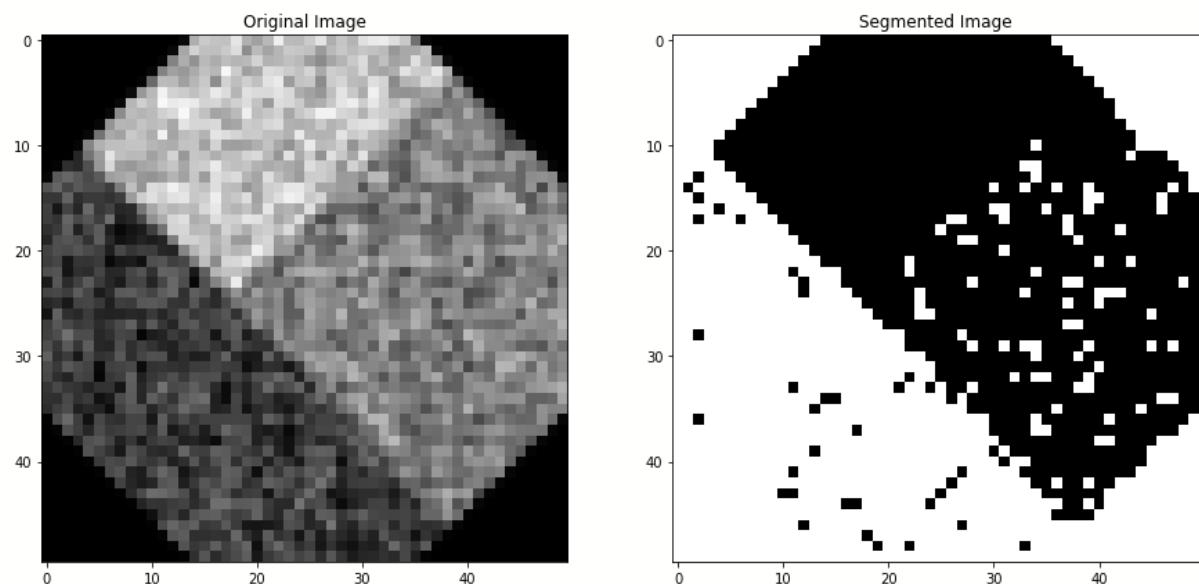


test5.jpg (downscaled - (50\*50) )

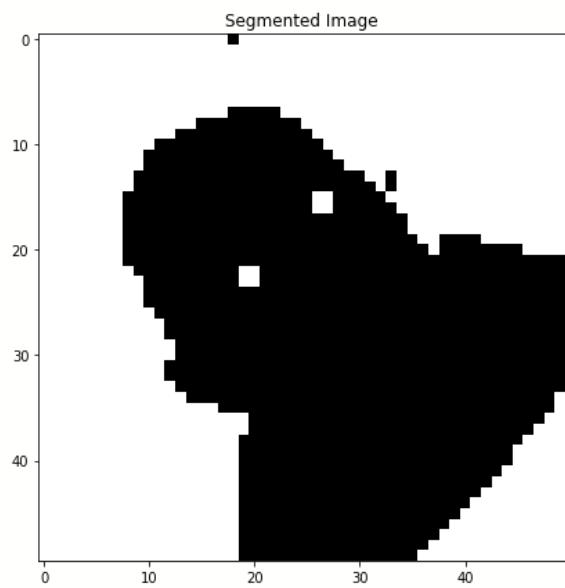
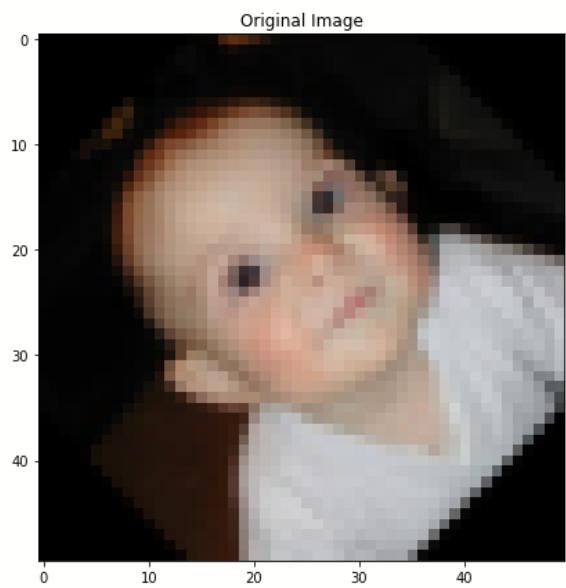


ii) N-Cut on Rotated Image (Downscaled)

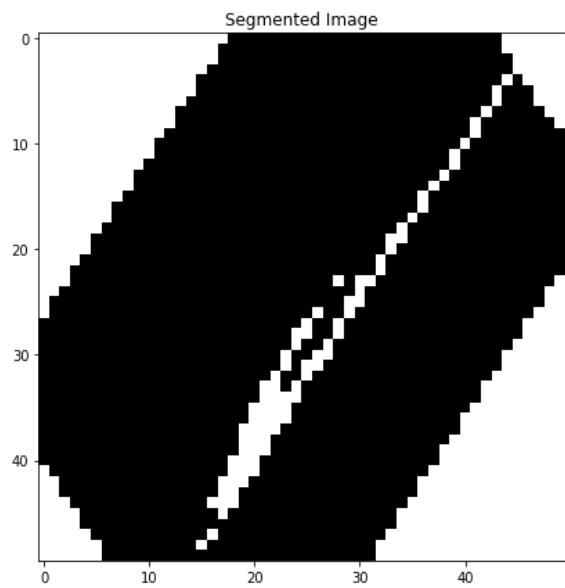
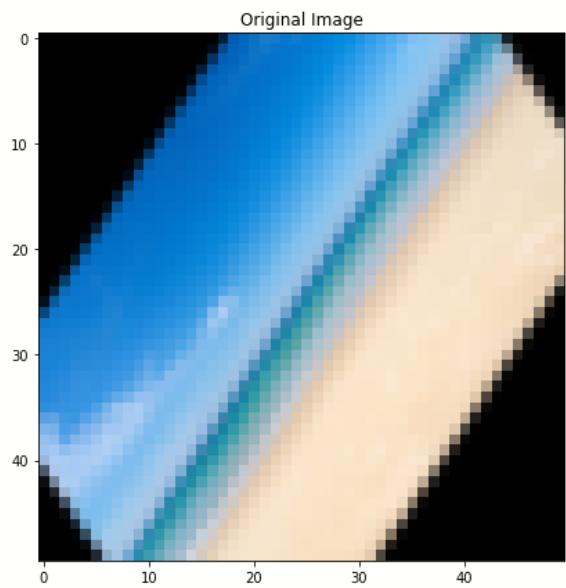
test1.png (downscaled - (50\*50) )



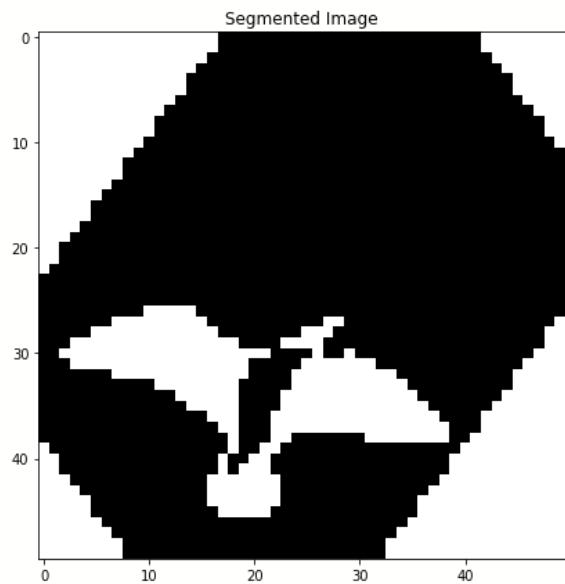
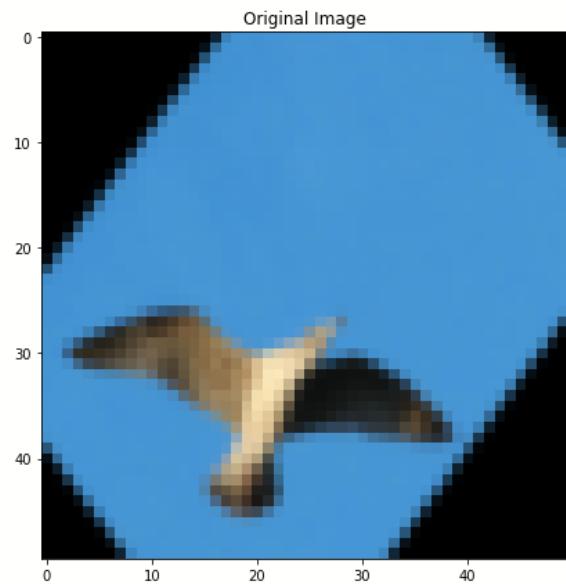
test2.jpg (downscaled - (50\*50) )



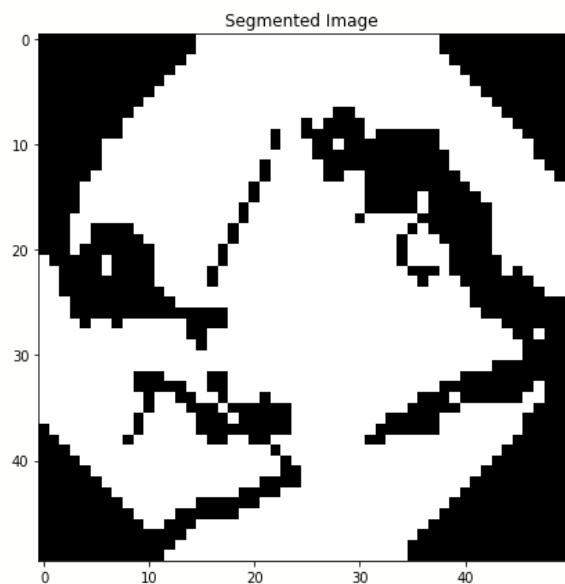
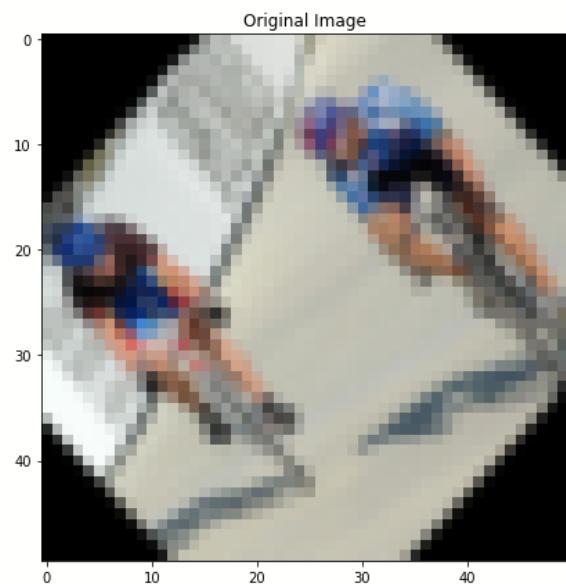
test3.jpg (downscaled - (50\*50) )



test4.jpg (downscaled - (50\*50) )

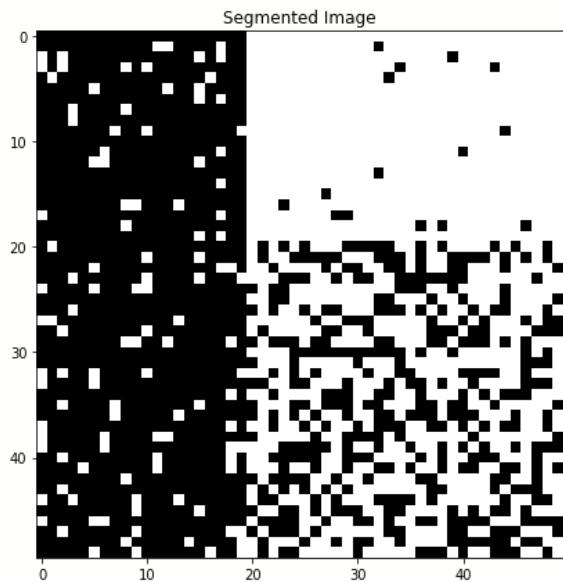
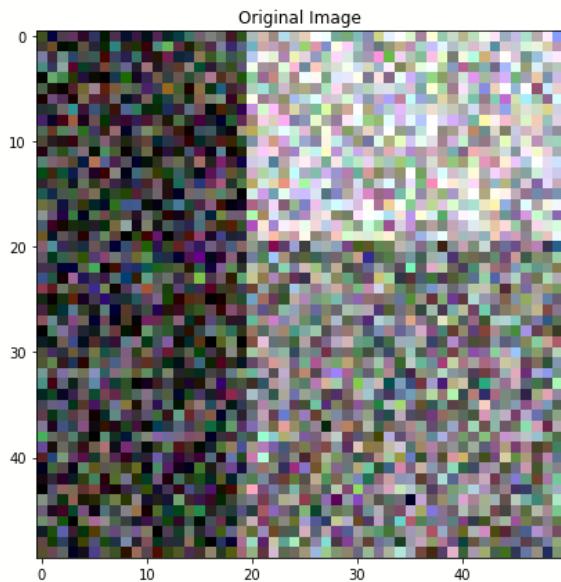


test5.jpg (downscaled - (50\*50) )

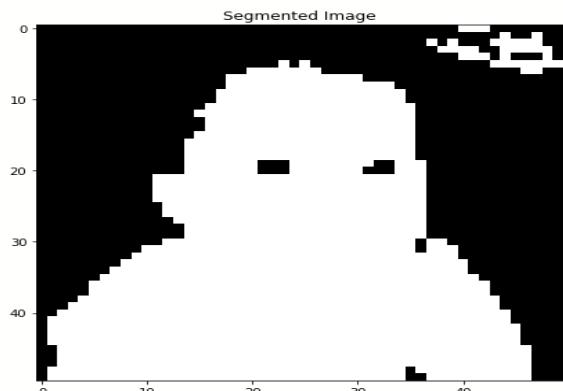


iii) N-Cut on Gaussian Blurred Image (Downscaled)

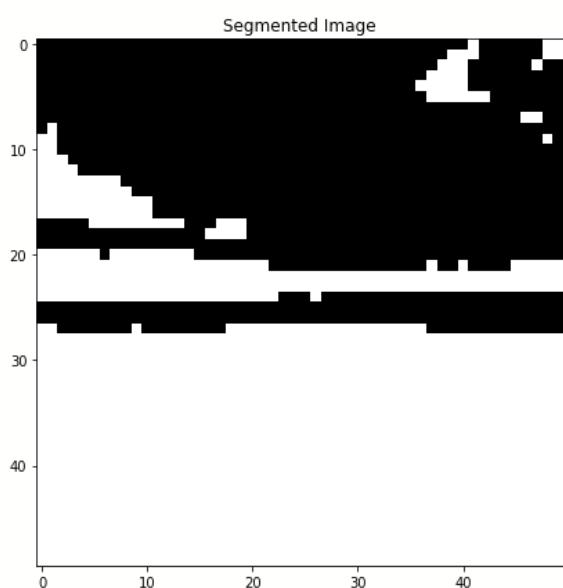
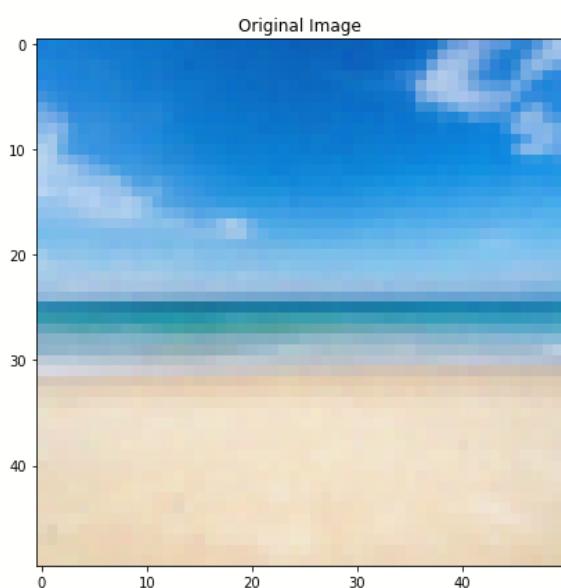
test1.png (downscaled - (50\*50) )



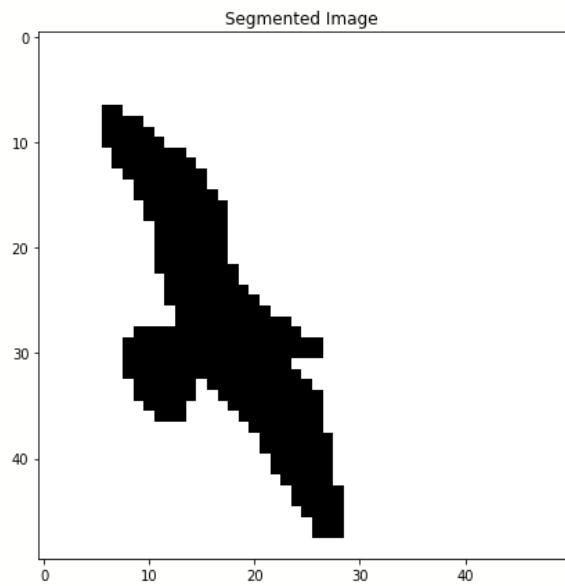
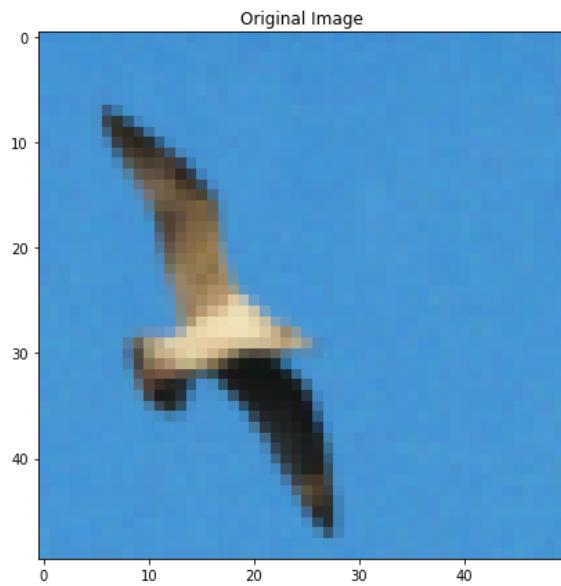
test2.jpg (downscaled - (50\*50) )



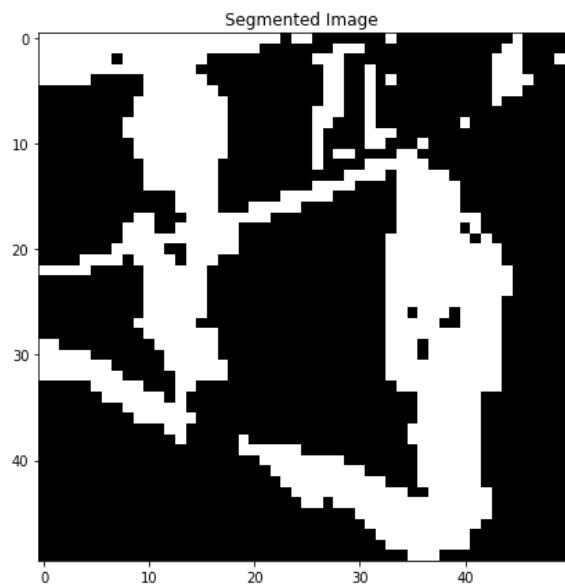
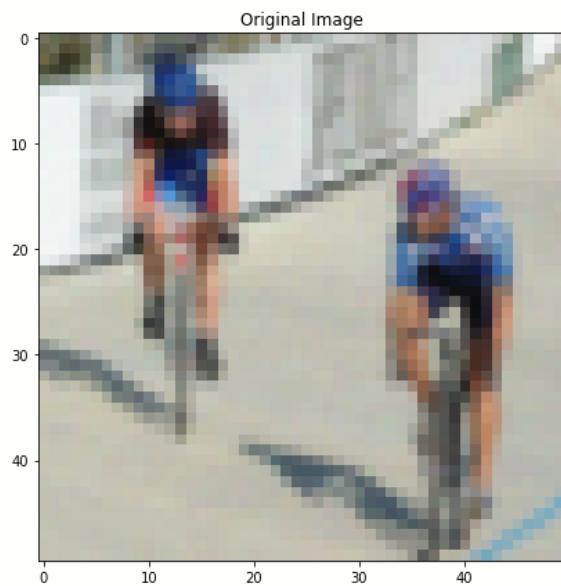
test3.jpg (downscaled - (50\*50) )



test4.jpg (downscaled - (50\*50) )



test5.jpg (downscaled - (50\*50) )



## Q2) Fully Convolutional Networks

- *Image segmentation using ResNet50 based FCN*
- Created a custom dataset class (`SegmentationDataset`) for the given dataset and implemented `__init__`, `len` and `get_item` methods in the class
- Split the trainval dataset into train and val data using the provided `train.txt` and `val.txt` entries
- Applied following data transformations on images and masks-  
Image - Transforms  
`transforms.Resize((520, 520)),`  
`transforms.ToTensor(),`  
`transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])`

```
Mask - Transforms
resize_transform =
transforms.Resize(520, 520), interpolation=InterpolationMode.NEAREST)
convert_tensor = transforms.PILToTensor()
```

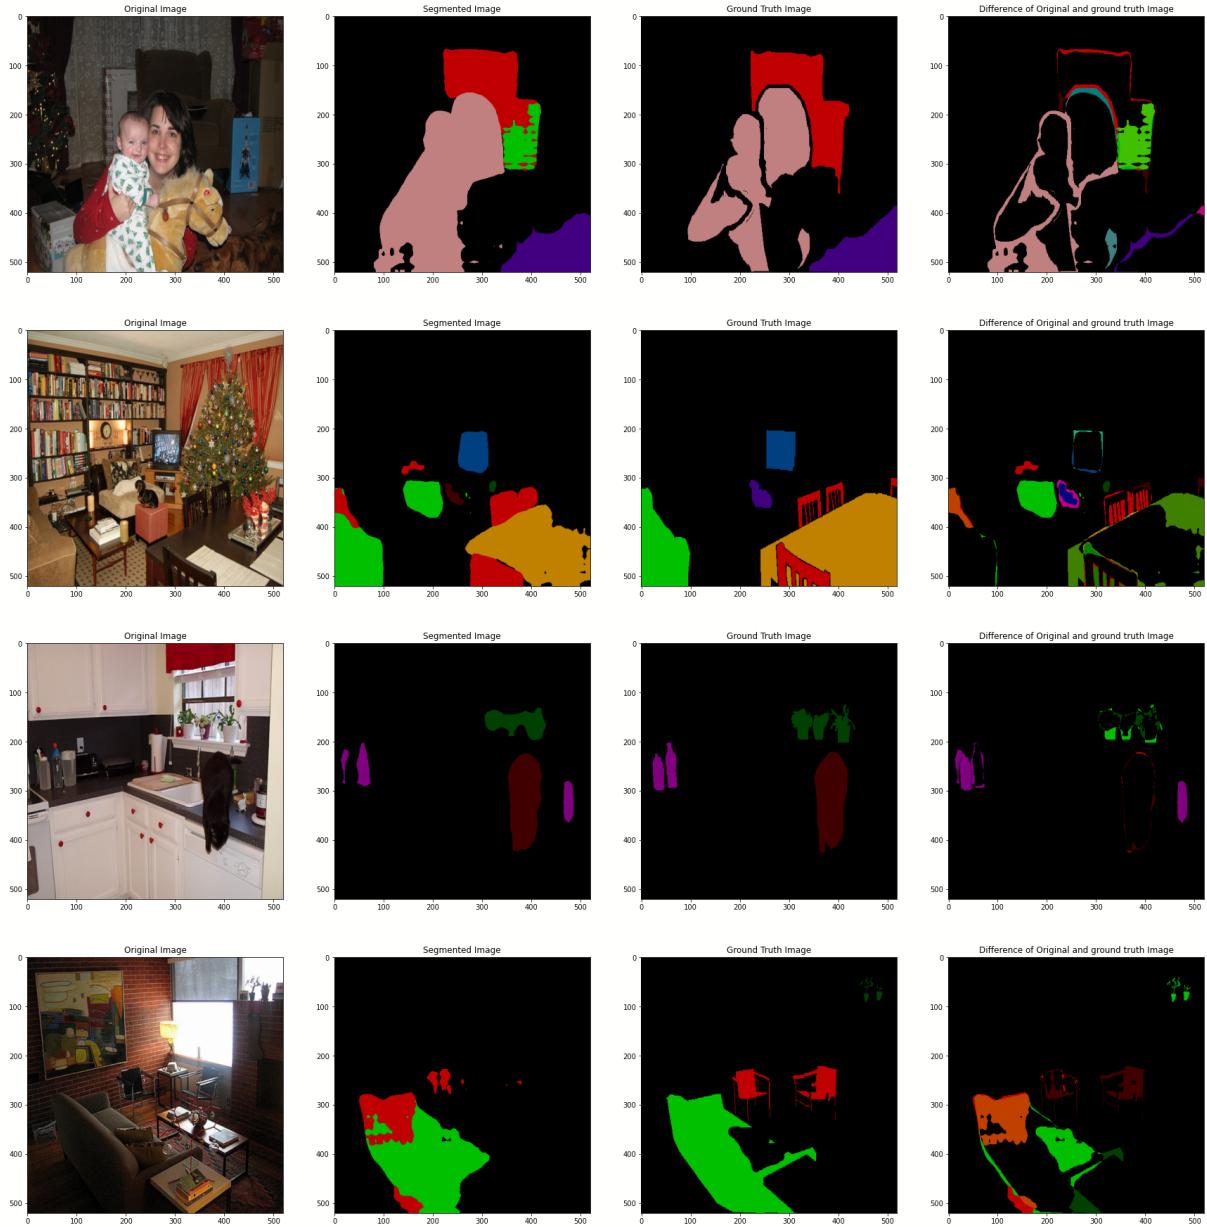
Note: `InterpolationMode.NEAREST` is used while resizing the image, so that colour corresponding to a label does not change in the mask.

- Loaded `fcn_resnet50` model with pretrained weights `FCN_ResNet50_Weights.DEFAULT` (weights obtained by training on COCO datasets)
- Tested the model on the provided test dataset and calculated the `Pixel_accuracy` and `MIOU`

### ● **Results:**

- Obtained the following results on the provided test dataset-
  - **Pixel Accuracy** - 87.92256797689491%
  - **MIOU** - 92.14278211168683%

- Visualisation of few outputs of model



- **Image Segmentation using FCN (MobileNetv2 backbone)**
- Created a custom dataset class (`SegmentationDataset`) for the given dataset and implemented `__init__`, `len` and `get_item` methods in the class
- Split the trainval dataset into train and val data using the provided `train.txt` and `val.txt` entries
- Applied following data transformations on images and masks-
  - Image - Transforms

```

transforms.Resize((520, 520)),
transforms.ToTensor(),
transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])

```

```

Mask - Transforms
resize_transform =
transforms.Resize((520, 520), interpolation=InterpolationMode.NEAREST)
convert_tensor = transforms.PILToTensor()

```

Note: InterpolationMode.NEAREST is used while resizing the image, so that colour corresponding to a label does not change in the mask.

- Created a FCN using MobileNetV2 as backbone , removing the final fc layer and then adding 5 upsample layers to get a final output of size [21, 224, 224] (21- classes, (224,224) - H\*W of input image)

Note-1: Each Upsample layer upscales the input by a factor of 2 using nn.ConvTranspose2d module.

Note-2: Final convolution layer of MobileNetV2 output size is 7\*7(h\*w) when input size is 224\*224(h\*w), so we upsample it 5 times.

- Weights of the nn.ConvTranspose2d are initialised using bilinear interpolation. (Observed that initialising weights in this way gives much better results compared to random initialization.)
- Model Summary-

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 112, 112]	864
BatchNorm2d-2	[-1, 32, 112, 112]	64
ReLU6-3	[-1, 32, 112, 112]	0
Conv2d-4	[-1, 32, 112, 112]	288
BatchNorm2d-5	[-1, 32, 112, 112]	64
ReLU6-6	[-1, 32, 112, 112]	0
Conv2d-7	[-1, 16, 112, 112]	512
BatchNorm2d-8	[-1, 16, 112, 112]	32
InvertedResidual-9	[-1, 16, 112, 112]	0
Conv2d-10	[-1, 96, 112, 112]	1,536
BatchNorm2d-11	[-1, 96, 112, 112]	192
ReLU6-12	[-1, 96, 112, 112]	0
Conv2d-13	[-1, 96, 56, 56]	864
BatchNorm2d-14	[-1, 96, 56, 56]	192

ReLU6-15	[-1, 96, 56, 56]	0
Conv2d-16	[-1, 24, 56, 56]	2,304
BatchNorm2d-17	[-1, 24, 56, 56]	48
InvertedResidual-18	[-1, 24, 56, 56]	0
Conv2d-19	[-1, 144, 56, 56]	3,456
BatchNorm2d-20	[-1, 144, 56, 56]	288
ReLU6-21	[-1, 144, 56, 56]	0
Conv2d-22	[-1, 144, 56, 56]	1,296
BatchNorm2d-23	[-1, 144, 56, 56]	288
ReLU6-24	[-1, 144, 56, 56]	0
Conv2d-25	[-1, 24, 56, 56]	3,456
BatchNorm2d-26	[-1, 24, 56, 56]	48
InvertedResidual-27	[-1, 24, 56, 56]	0
Conv2d-28	[-1, 144, 56, 56]	3,456
BatchNorm2d-29	[-1, 144, 56, 56]	288
ReLU6-30	[-1, 144, 56, 56]	0
Conv2d-31	[-1, 144, 28, 28]	1,296
BatchNorm2d-32	[-1, 144, 28, 28]	288
ReLU6-33	[-1, 144, 28, 28]	0
Conv2d-34	[-1, 32, 28, 28]	4,608
BatchNorm2d-35	[-1, 32, 28, 28]	64
InvertedResidual-36	[-1, 32, 28, 28]	0
Conv2d-37	[-1, 192, 28, 28]	6,144
BatchNorm2d-38	[-1, 192, 28, 28]	384
ReLU6-39	[-1, 192, 28, 28]	0
Conv2d-40	[-1, 192, 28, 28]	1,728
BatchNorm2d-41	[-1, 192, 28, 28]	384
ReLU6-42	[-1, 192, 28, 28]	0
Conv2d-43	[-1, 32, 28, 28]	6,144
BatchNorm2d-44	[-1, 32, 28, 28]	64
InvertedResidual-45	[-1, 32, 28, 28]	0
Conv2d-46	[-1, 192, 28, 28]	6,144
BatchNorm2d-47	[-1, 192, 28, 28]	384
ReLU6-48	[-1, 192, 28, 28]	0
Conv2d-49	[-1, 192, 28, 28]	1,728
BatchNorm2d-50	[-1, 192, 28, 28]	384
ReLU6-51	[-1, 192, 28, 28]	0
Conv2d-52	[-1, 32, 28, 28]	6,144
BatchNorm2d-53	[-1, 32, 28, 28]	64
InvertedResidual-54	[-1, 32, 28, 28]	0
Conv2d-55	[-1, 192, 28, 28]	6,144
BatchNorm2d-56	[-1, 192, 28, 28]	384
ReLU6-57	[-1, 192, 28, 28]	0

Conv2d-58	$[-1, 192, 14, 14]$	1,728
BatchNorm2d-59	$[-1, 192, 14, 14]$	384
ReLU6-60	$[-1, 192, 14, 14]$	0
Conv2d-61	$[-1, 64, 14, 14]$	12,288
BatchNorm2d-62	$[-1, 64, 14, 14]$	128
InvertedResidual-63	$[-1, 64, 14, 14]$	0
Conv2d-64	$[-1, 384, 14, 14]$	24,576
BatchNorm2d-65	$[-1, 384, 14, 14]$	768
ReLU6-66	$[-1, 384, 14, 14]$	0
Conv2d-67	$[-1, 384, 14, 14]$	3,456
BatchNorm2d-68	$[-1, 384, 14, 14]$	768
ReLU6-69	$[-1, 384, 14, 14]$	0
Conv2d-70	$[-1, 64, 14, 14]$	24,576
BatchNorm2d-71	$[-1, 64, 14, 14]$	128
InvertedResidual-72	$[-1, 64, 14, 14]$	0
Conv2d-73	$[-1, 384, 14, 14]$	24,576
BatchNorm2d-74	$[-1, 384, 14, 14]$	768
ReLU6-75	$[-1, 384, 14, 14]$	0
Conv2d-76	$[-1, 384, 14, 14]$	3,456
BatchNorm2d-77	$[-1, 384, 14, 14]$	768
ReLU6-78	$[-1, 384, 14, 14]$	0
Conv2d-79	$[-1, 64, 14, 14]$	24,576
BatchNorm2d-80	$[-1, 64, 14, 14]$	128
InvertedResidual-81	$[-1, 64, 14, 14]$	0
Conv2d-82	$[-1, 384, 14, 14]$	24,576
BatchNorm2d-83	$[-1, 384, 14, 14]$	768
ReLU6-84	$[-1, 384, 14, 14]$	0
Conv2d-85	$[-1, 384, 14, 14]$	3,456
BatchNorm2d-86	$[-1, 384, 14, 14]$	768
ReLU6-87	$[-1, 384, 14, 14]$	0
Conv2d-88	$[-1, 64, 14, 14]$	24,576
BatchNorm2d-89	$[-1, 64, 14, 14]$	128
InvertedResidual-90	$[-1, 64, 14, 14]$	0
Conv2d-91	$[-1, 384, 14, 14]$	24,576
BatchNorm2d-92	$[-1, 384, 14, 14]$	768
ReLU6-93	$[-1, 384, 14, 14]$	0
Conv2d-94	$[-1, 384, 14, 14]$	3,456
BatchNorm2d-95	$[-1, 384, 14, 14]$	768
ReLU6-96	$[-1, 384, 14, 14]$	0
Conv2d-97	$[-1, 96, 14, 14]$	36,864
BatchNorm2d-98	$[-1, 96, 14, 14]$	192
InvertedResidual-99	$[-1, 96, 14, 14]$	0
Conv2d-100	$[-1, 576, 14, 14]$	55,296

BatchNorm2d-101	$[-1, 576, 14, 14]$	1,152
ReLU6-102	$[-1, 576, 14, 14]$	0
Conv2d-103	$[-1, 576, 14, 14]$	5,184
BatchNorm2d-104	$[-1, 576, 14, 14]$	1,152
ReLU6-105	$[-1, 576, 14, 14]$	0
Conv2d-106	$[-1, 96, 14, 14]$	55,296
BatchNorm2d-107	$[-1, 96, 14, 14]$	192
InvertedResidual-108	$[-1, 96, 14, 14]$	0
Conv2d-109	$[-1, 576, 14, 14]$	55,296
BatchNorm2d-110	$[-1, 576, 14, 14]$	1,152
ReLU6-111	$[-1, 576, 14, 14]$	0
Conv2d-112	$[-1, 576, 14, 14]$	5,184
BatchNorm2d-113	$[-1, 576, 14, 14]$	1,152
ReLU6-114	$[-1, 576, 14, 14]$	0
Conv2d-115	$[-1, 96, 14, 14]$	55,296
BatchNorm2d-116	$[-1, 96, 14, 14]$	192
InvertedResidual-117	$[-1, 96, 14, 14]$	0
Conv2d-118	$[-1, 576, 14, 14]$	55,296
BatchNorm2d-119	$[-1, 576, 14, 14]$	1,152
ReLU6-120	$[-1, 576, 14, 14]$	0
Conv2d-121	$[-1, 576, 7, 7]$	5,184
BatchNorm2d-122	$[-1, 576, 7, 7]$	1,152
ReLU6-123	$[-1, 576, 7, 7]$	0
Conv2d-124	$[-1, 160, 7, 7]$	92,160
BatchNorm2d-125	$[-1, 160, 7, 7]$	320
InvertedResidual-126	$[-1, 160, 7, 7]$	0
Conv2d-127	$[-1, 960, 7, 7]$	153,600
BatchNorm2d-128	$[-1, 960, 7, 7]$	1,920
ReLU6-129	$[-1, 960, 7, 7]$	0
Conv2d-130	$[-1, 960, 7, 7]$	8,640
BatchNorm2d-131	$[-1, 960, 7, 7]$	1,920
ReLU6-132	$[-1, 960, 7, 7]$	0
Conv2d-133	$[-1, 160, 7, 7]$	153,600
BatchNorm2d-134	$[-1, 160, 7, 7]$	320
InvertedResidual-135	$[-1, 160, 7, 7]$	0
Conv2d-136	$[-1, 960, 7, 7]$	153,600
BatchNorm2d-137	$[-1, 960, 7, 7]$	1,920
ReLU6-138	$[-1, 960, 7, 7]$	0
Conv2d-139	$[-1, 960, 7, 7]$	8,640
BatchNorm2d-140	$[-1, 960, 7, 7]$	1,920
ReLU6-141	$[-1, 960, 7, 7]$	0
Conv2d-142	$[-1, 160, 7, 7]$	153,600
BatchNorm2d-143	$[-1, 160, 7, 7]$	320

InvertedResidual-144	$[-1, 160, 7, 7]$	0
Conv2d-145	$[-1, 960, 7, 7]$	153,600
BatchNorm2d-146	$[-1, 960, 7, 7]$	1,920
ReLU6-147	$[-1, 960, 7, 7]$	0
Conv2d-148	$[-1, 960, 7, 7]$	8,640
BatchNorm2d-149	$[-1, 960, 7, 7]$	1,920
ReLU6-150	$[-1, 960, 7, 7]$	0
Conv2d-151	$[-1, 320, 7, 7]$	307,200
BatchNorm2d-152	$[-1, 320, 7, 7]$	640
InvertedResidual-153	$[-1, 320, 7, 7]$	0
Conv2d-154	$[-1, 1280, 7, 7]$	409,600
BatchNorm2d-155	$[-1, 1280, 7, 7]$	2,560
ReLU6-156	$[-1, 1280, 7, 7]$	0
<b>ConvTranspose2d-157</b>	<b><math>[-1, 512, 14, 14]</math></b>	<b>10,486,272</b>
BatchNorm2d-158	$[-1, 512, 14, 14]$	1,024
ReLU-159	$[-1, 512, 14, 14]$	0
UpsampleBlock-160	$[-1, 512, 14, 14]$	0
<b>ConvTranspose2d-161</b>	<b><math>[-1, 256, 28, 28]</math></b>	<b>2,097,408</b>
BatchNorm2d-162	$[-1, 256, 28, 28]$	512
ReLU-163	$[-1, 256, 28, 28]$	0
UpsampleBlock-164	$[-1, 256, 28, 28]$	0
<b>ConvTranspose2d-165</b>	<b><math>[-1, 128, 56, 56]</math></b>	<b>524,416</b>
BatchNorm2d-166	$[-1, 128, 56, 56]$	256
ReLU-167	$[-1, 128, 56, 56]$	0
UpsampleBlock-168	$[-1, 128, 56, 56]$	0
<b>ConvTranspose2d-169</b>	<b><math>[-1, 64, 112, 112]</math></b>	<b>131,136</b>
BatchNorm2d-170	$[-1, 64, 112, 112]$	128
ReLU-171	$[-1, 64, 112, 112]$	0
UpsampleBlock-172	$[-1, 64, 112, 112]$	0
<b>ConvTranspose2d-173</b>	<b><math>[-1, 21, 224, 224]</math></b>	<b>21,525</b>
BatchNorm2d-174	$[-1, 21, 224, 224]$	42
ReLU-175	$[-1, 21, 224, 224]$	0
UpsampleBlock-176	$[-1, 21, 224, 224]$	0

=====

Total params: 15,486,591

Trainable params: 13,262,719

Non-trainable params: 2,223,872

-----

Input size (MB): 0.57

Forward/backward pass size (MB): 230.94

Params size (MB): 59.08

Estimated Total Size (MB): 290.59

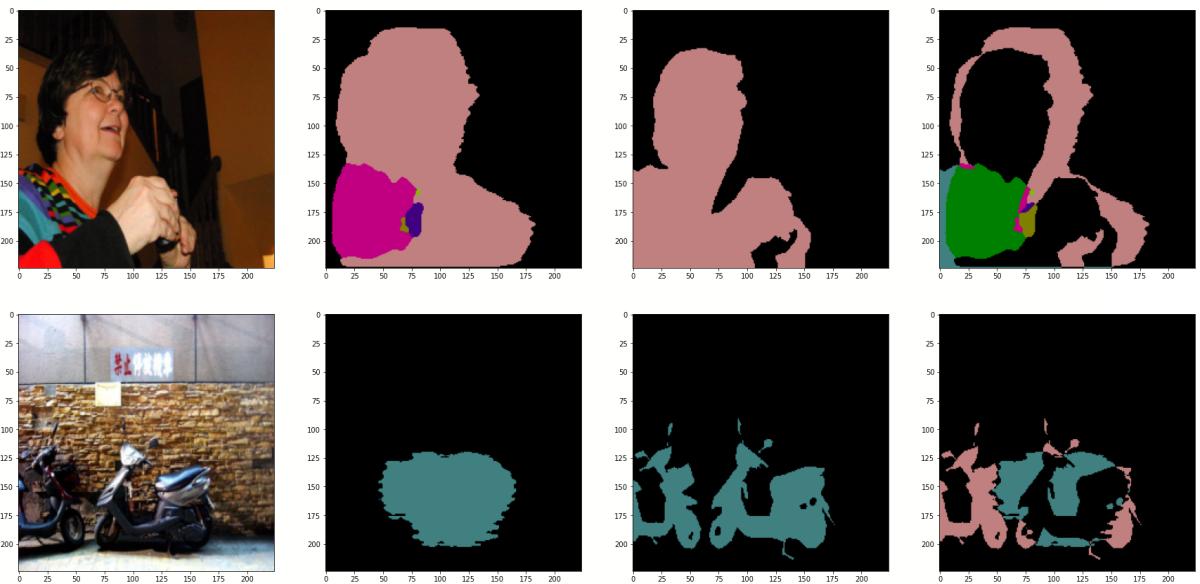
-----

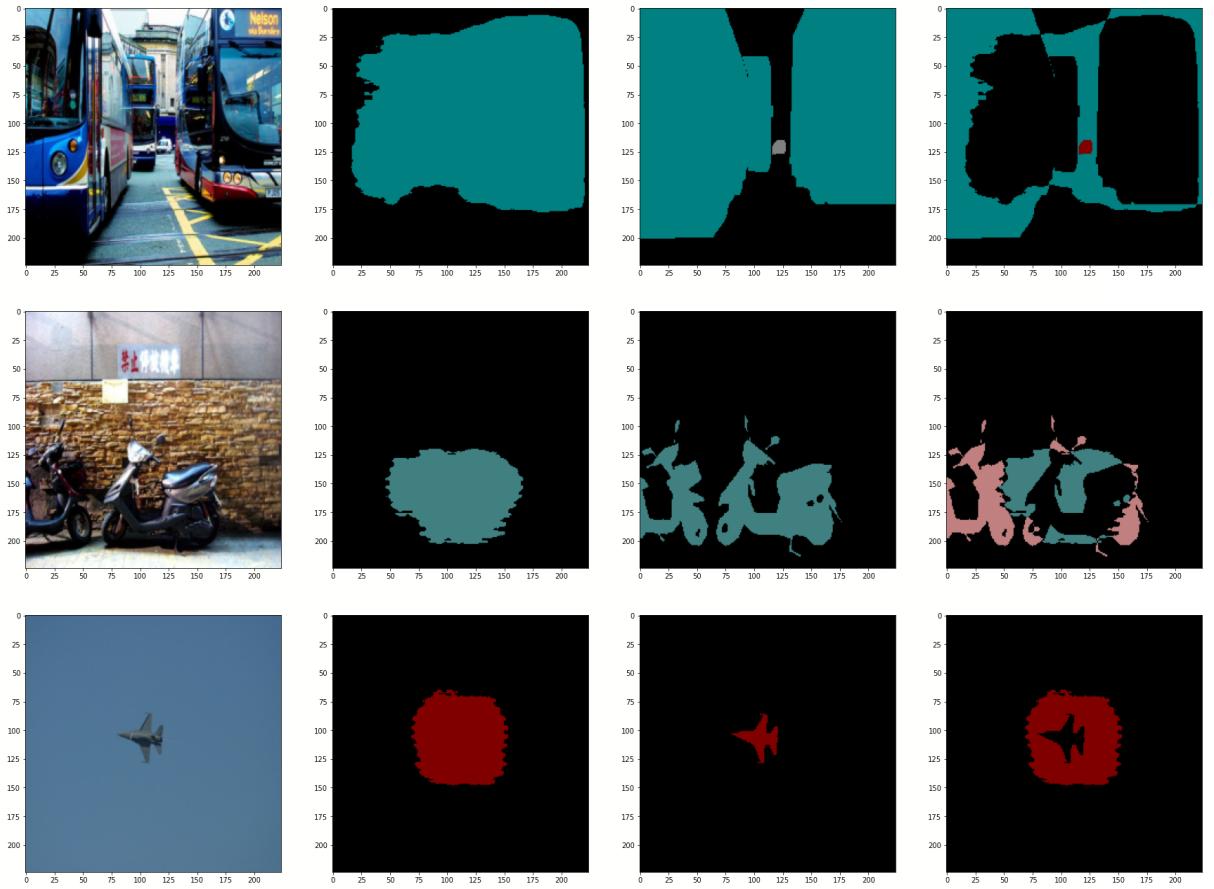
Note: Layers added by me are in **bold**

- Trained the model on provided training data using the following hyperparameters
  - `optimizer_ft = optim.SGD(params_to_update, lr=0.001, momentum=0.9)`
  - `criterion = nn.CrossEntropyLoss(ignore_index=255)`
  - `exp_lr_scheduler = lr_scheduler.StepLR(optimizer_ft, step_size=70, gamma=1)`
  - Epochs - 150
- Tested the model on the provided test dataset and calculated the pixel accuracy and MIOU.

### ● **Results:**

- Obtained the following results on the provided test dataset-
  - **Pixel Accuracy** - 71.20890655369291%
  - **MIOU** - 81.37229567087272%
- Visualisation of few outputs of model-





Note: 1st Image- Original Image

2nd Image- Model Output

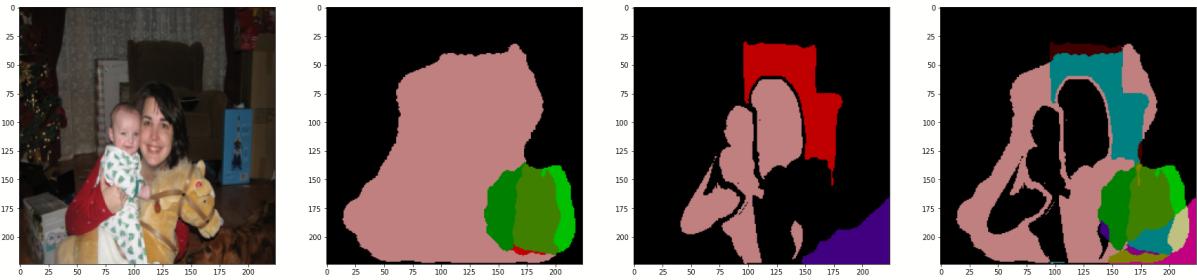
3rd Image- Ground Truth

4th Image- Difference between Model Output and ground truth

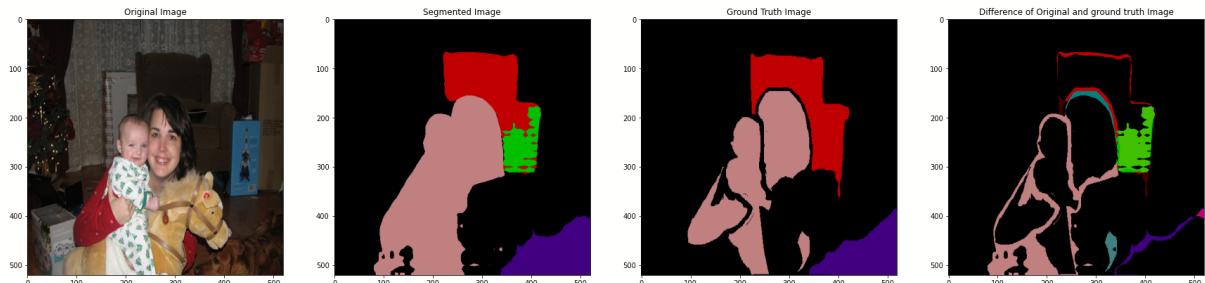
### ● ***Comparison Between the two Models***

- From the calculated model accuracies and visualised outputs, it is clear that FCN ResNet50 with pretrained weights is performing better than trained MobileNetV2 backbone FCN.
- Lower performance of trained MobileNetV2 backbone FCN can attributed to the following reasons
  - No Skip connections in the model (Adding skip connections will lead to much better performance of our model)
  - Insufficient data- ResNet50 FCN is trained on COCO Dataset which is very large compared to PASCAL VOC dataset (a subset of it)

- Comparison of few visualised outputs



Output of MobileNetV2 backbone FCN



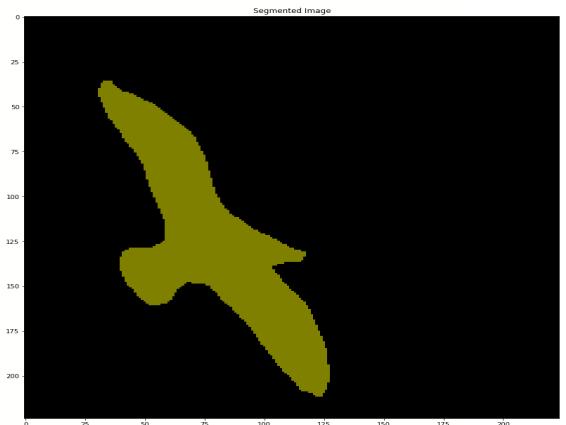
Output of ResNet50 FCN

- ***Comparison between N-Cut and FCN's***
- Compared N-cut and FCN's by segmenting test4.jpg and test5.jpg images using both N-cut and FCN's
- Observed that the segmented image by N-cut and ResNet50 FCN are good, but computation time for N-cut is much higher than ResNet50 FCN.
- MobileNetV2 backbone FCN did not perform well on these images due to the reasons mentioned in the previous section.

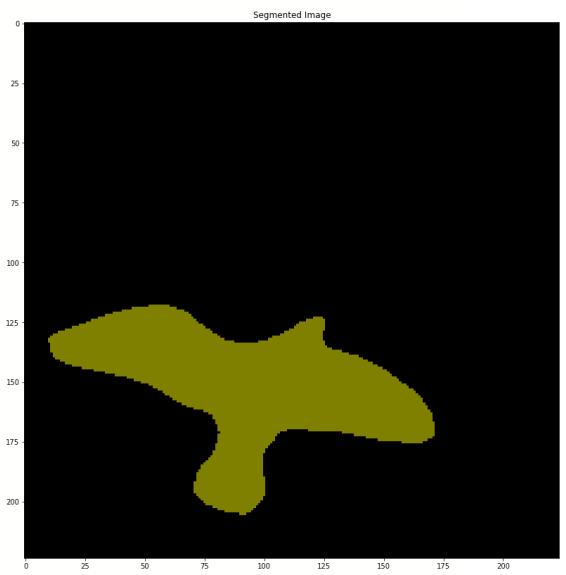
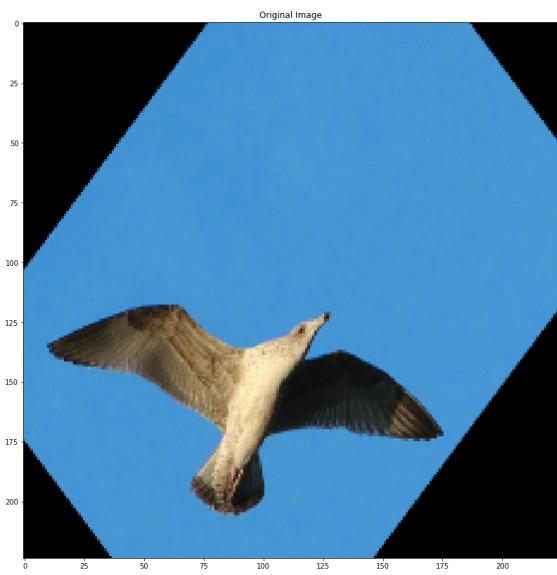
- Visualisation of segmented images by FCN's

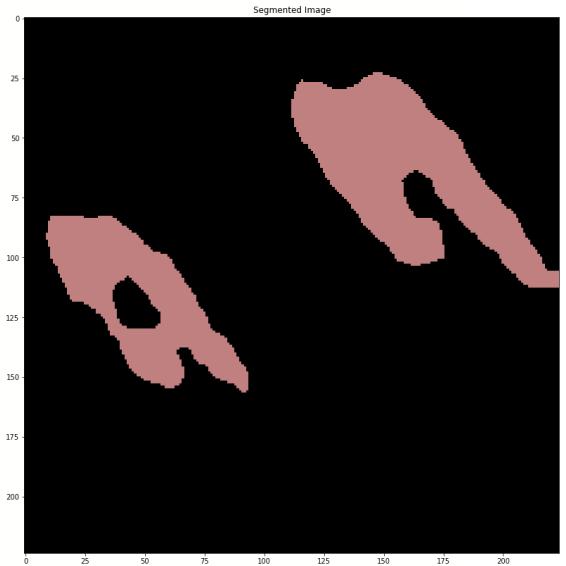
## Outputs of ResNet-50 FCN

### Original Images

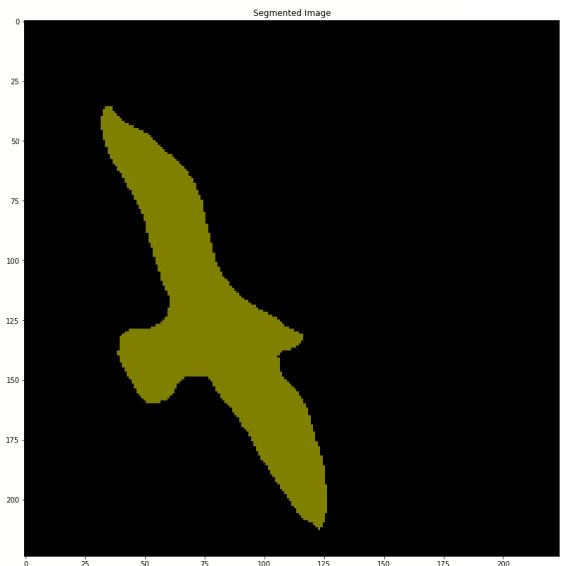
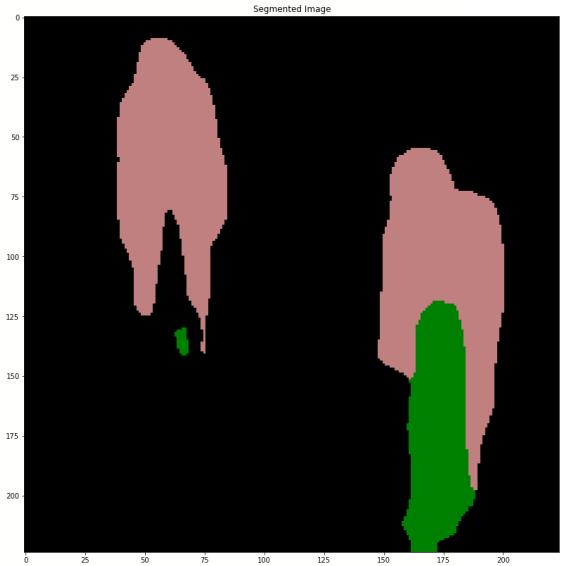
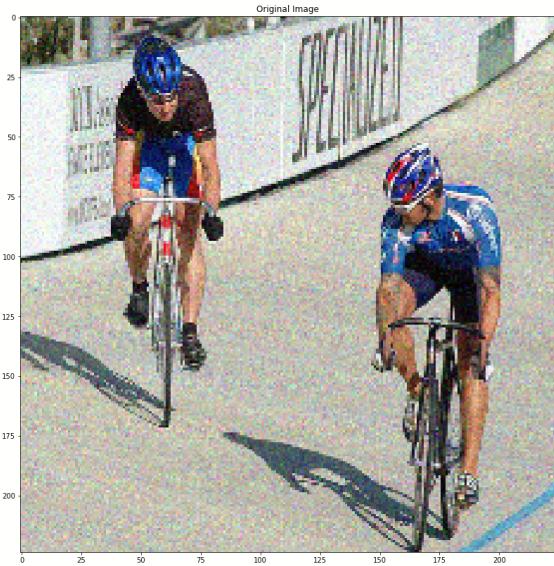


### Rotated Images



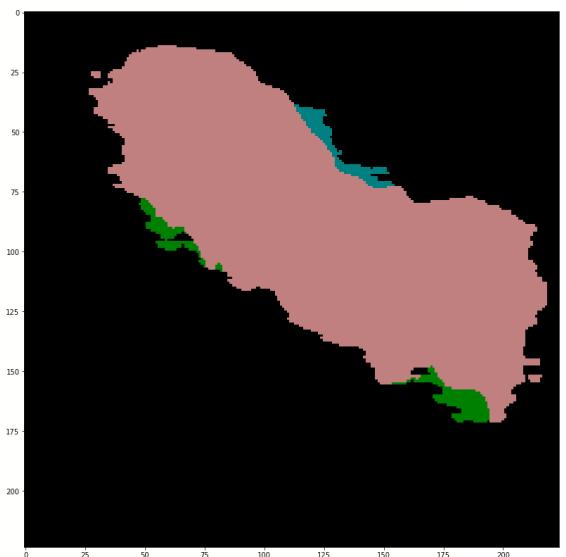
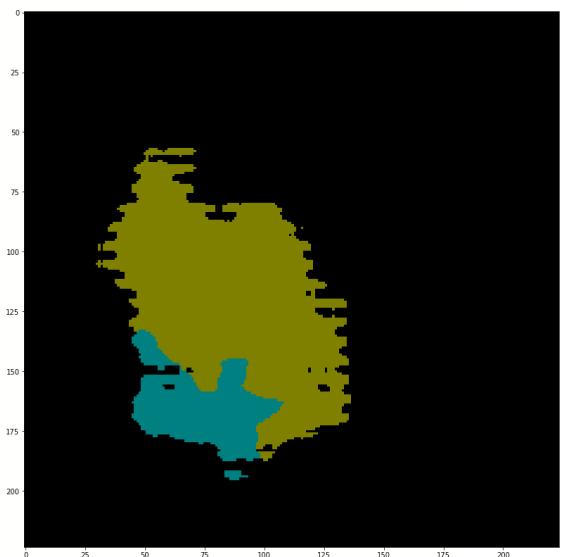


Noisy Images

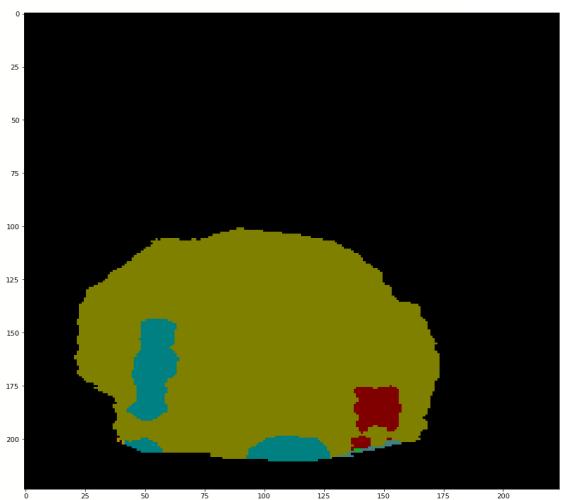


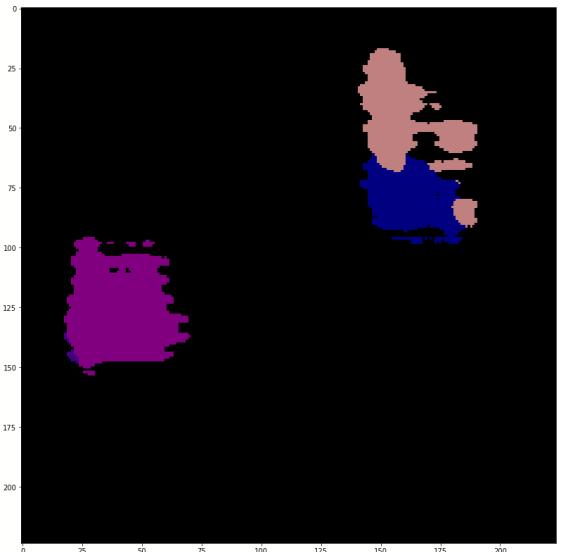
### Outputs of MobileNetV2 Backbone FCN

#### Original Images

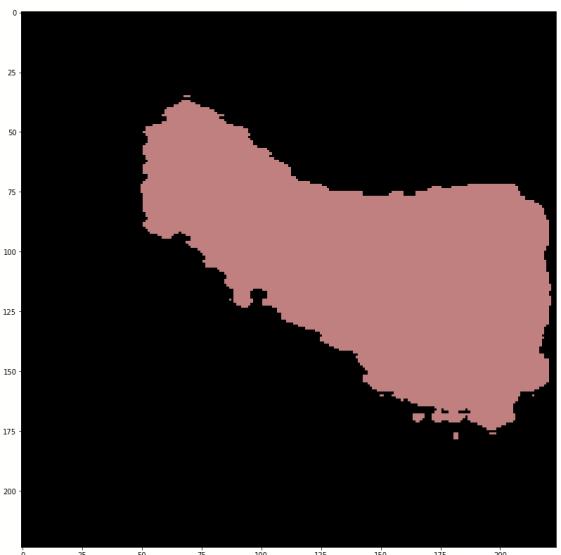
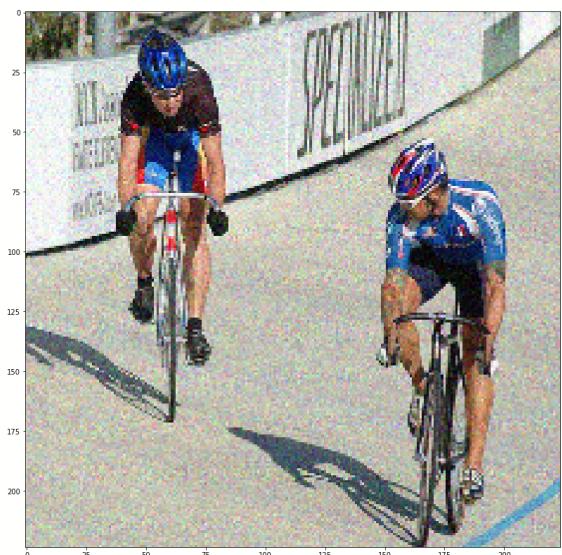
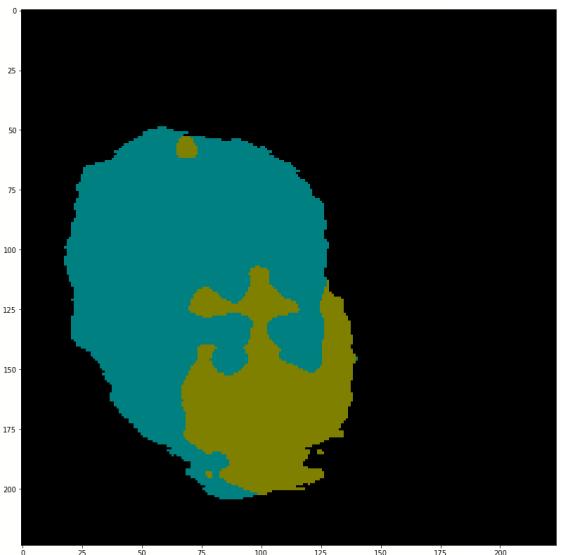
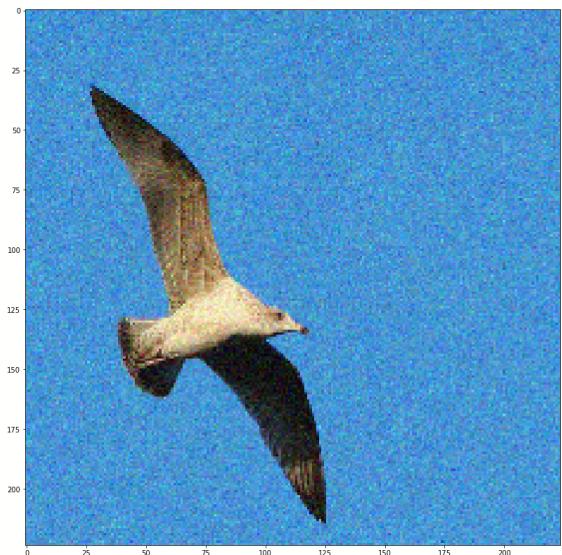


#### Rotated Images





Noisy Images



**References:**

[https://d2l.ai/chapter\\_computer-vision/fcn.html](https://d2l.ai/chapter_computer-vision/fcn.html)  
[https://github.com/sagieppel/Fully-convolutional-neural-network-FCN-for-semantic-segmentation-with-pytorch/blob/master/NET\\_FCN.py](https://github.com/sagieppel/Fully-convolutional-neural-network-FCN-for-semantic-segmentation-with-pytorch/blob/master/NET_FCN.py)  
[https://poutyne.org/examples/semantic\\_segmentation.html](https://poutyne.org/examples/semantic_segmentation.html)  
[https://pytorch.org/vision/main/\\_modules/torchvision/datasets/voc.html#VOCSegmentation](https://pytorch.org/vision/main/_modules/torchvision/datasets/voc.html#VOCSegmentation)